


# Week 9

## Agenda:

1. Vector Processors and RISC-V Vector Extension
2. SiFive Vector Processor


# Vector Processors

**SCALAR**  
(1 operation)



add r3, r1, r2

**VECTOR**  
(N operations)




vadd.vv v3, v1, v2

Source : Patterson Lecture , Vector Processor

## Constituents of Vector Processor

1. **Vector Register File:** Each vector register is fixed length bank holding a single vector
  - Has atleast 2 read and 1 write ports
  - 32 Vector Regs in RISC-V – V Spec
2. **Vector Functional Unit(FU):** Fully pipelined to start new operation every clock. Control unit to manage hazards
3. **Load Store Unit(LSU):** Fully pipelined unit to amortize the memory latency over one vector load.
4. **Crossbars** to connect LSU and FU



Source : Patterson Hennesy, Comp Arch Text



# RISC-V Foundation Vector Extension Overview

32 vector registers

v31[0]	v31[1]		v31[VLMAX-1]
v1[0]	v1[1]		v1[VLMAX-1]
v0[0]	v0[1]		v0[VLMAX-1]

Maximum vector length (VLMAX) depends on implementation, number of vector registers used, and type of each element.

- Unit-stride, strided, scatter-gather, structure load/store instructions
- Rich set of integer, fixed-point, and floating-point instructions
- Vector-vector, vector-scalar, and vector-immediate instructions
- Multiple vector registers can be combined to form longer vectors to reduce instruction bandwidth or support mixed-precision operations (e.g., 16b\*16b->32b multiply-accumulate)
- Designed for extension with custom datatypes and widths

Vector CSRs

vtype

Vtype sets width of element in each vector register (e.g., 16-bit, 32-bit, ...)

vl

Vector length CSR sets number of elements active in each instruction

vstart

Resumption element after trap

fcsr (vxrm/vxsat)

Fixed-point rounding mode and saturation flag fields in FP CSR

# RISC V Vector ISA

- 32 vector registers (v0 ... v31)
- Each register can hold either a scalar, a vector or a matrix(shape)
- Each vector register can optionally have an associated type(polymorphic encoding)
- Variable number of registers (dynamically changeable)
- **Vector instruction semantics**
- All instructions controlled by Vector Length (VL) register
- All instructions can be executed under mask
- Intuitive memory ordering model
- Precise exceptions supported
- **Vector instruction set:**
- All instructions present in base line ISA are present in the vector ISA
- Vector memory instructions supporting linear, strided & gather/scatter access patterns
- Optional Fixed-Point set
- Optional Transcendental set

# Week 10: Oct 2<sup>nd</sup> 2020

## **Actions**

1. Study RISC-V Vector Extension
2. Case Study RISCV Vector Processors

## **Work Done**

1. Study RISC-V Vector Extension Spec
2. Detailed study of Ara RISC-V Vector Processor

# Ara : RISC-V Vector Processor

## Overview

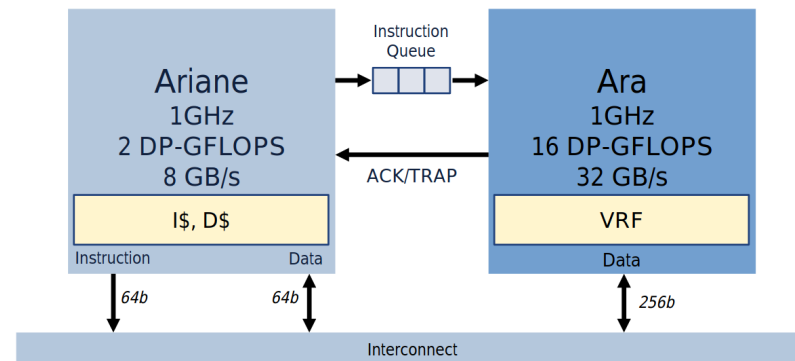
- Developed by PULP team, Integrated Systems Laboratory, ETH-Zurich in 2019
- 64-bit vector co-processor to Ariane (which is RV64GC application processor)
- Based on RISC-V Vector Spec. Version 0.5
- Operating Frequency 1 Ghz
- Designed for low power
- Contains Floating Point Support with 64 bit/cycle throughput (1 DP / 2 SP / 4 HP outputs/cycle)
- Peak Performance : 33 DP – GFLOPS (41 DP-GFLOPS/W)
- Exploits Data Level parallelism for efficient execution of Scientific and Matrix oriented computations, as well as DSP and Machine learning algorithms

# Topics of discussion

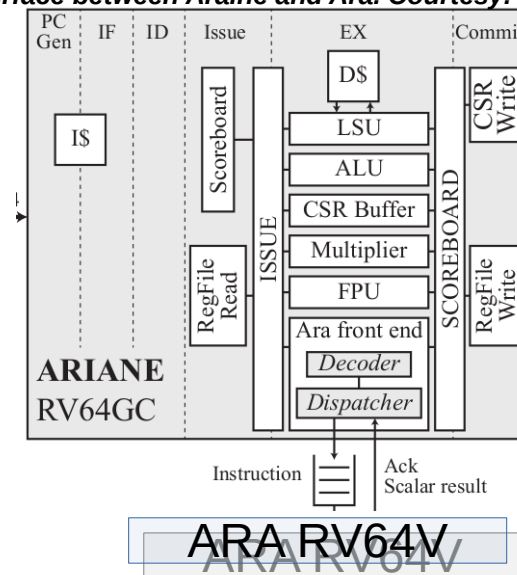
- ARA Microarchitecture
- ARA Performance Benchmarking

# Ariane Vector Processor System

- Ariane is a RV64GC based in-order , single issue application class processor having a 6-stage pipeline
- Vector instructions are partially decoded in Ariane's Instruction Decoder, and then completely in a dedicated Vector Instruction Decoder insider Ara.
- The dispatcher controls the interface between Ara and Ariane's dedicated Scoreboard port.
- While Ariane sends Instructions, Ara sends Ack/Scalar results through the Scoreboard port.
- The decoupled execution works well, except when Ariane expects a result from Ara, e.g., reading an element of a vector register



*Block Level Interface between Ariane and Ara. Courtesy: ETH Zurich ppt*

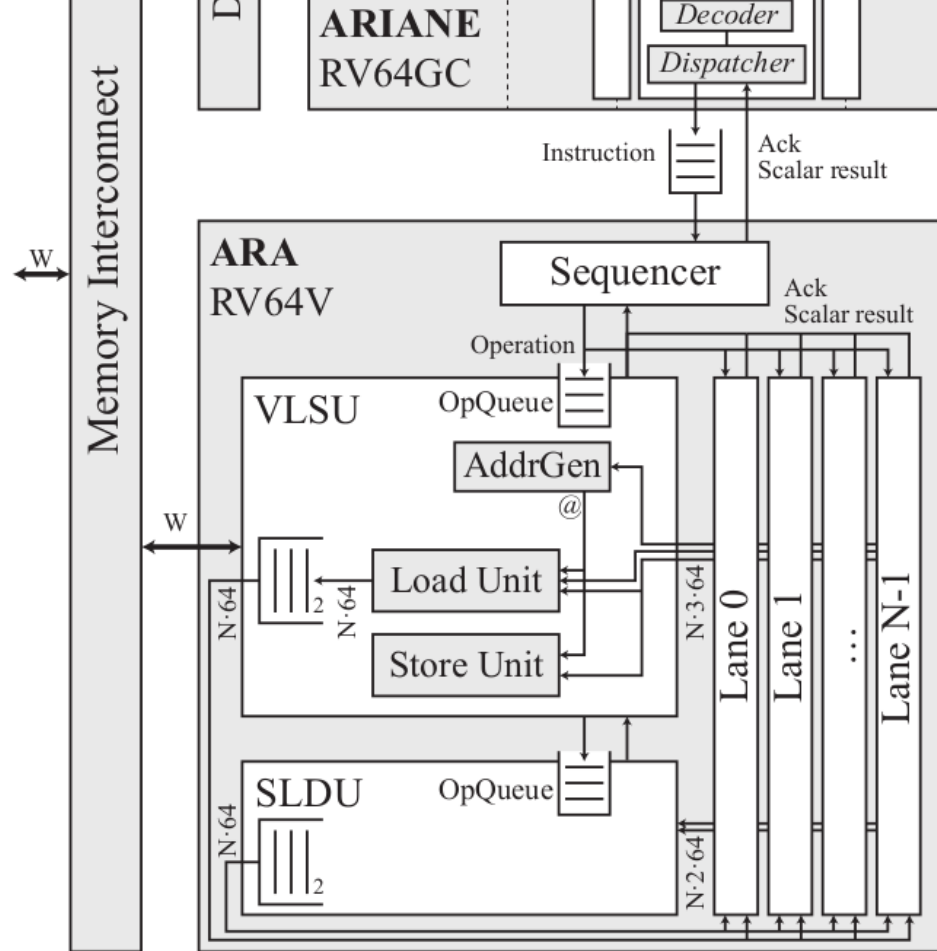


*Fine Grain Block Level Interface between Ariane and Ara. Courtesy: Ara Paper*



# Components of Ara

- Sequencer
- Slide Unit (SLDU)
- Vector Load Store unit (VLSU)
- Vector Lanes
  - Lane Sequencer
  - Vector Register File
  - Operand Queues
  - Execution Units



(a) Block diagram of an Ara instance with  $N$  parallel lanes. Ara receives its commands from Ariane, a RV64GC scalar core. The vector unit has a main sequencer;  $N$  parallel lanes; a Slide Unit (SLDU); and a Vector Load/Store Unit (VLSU). The memory interface is  $W$  bit wide.

Fig. 3. Top-level block diagram of Ara.

# Sequencer

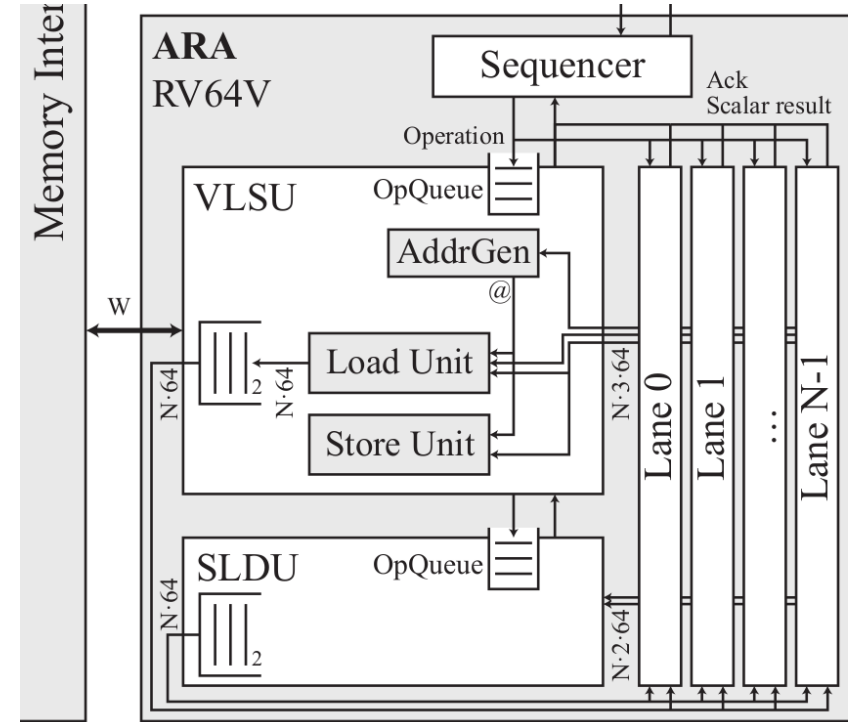
- Responsible for keeping track of vector instructions running on Ara, dispatching them to different execution units and acknowledging them with Ariane.
- This unit has a global view of the instruction execution progress across all lanes.
- The sequencer can handle upto 8 parallel instructions
- Hazards among pending vector instructions are resolved by this block.
- In case of a structural hazard, the sequencer delays the issue of vector instructions until the contending instruction releases the shared resource.
- Sequencer stores information about which vector instruction is accessing which vector register. This information is used to resolve data hazards between instructions

# Slide Unit (SLDU)

- Vector Register File (VRF) is spread across all lanes. Thus, a given lane cannot access all VRF banks directly.
- SLDU is responsible for handling instructions that must access all Vector Register File (VRF) banks at once.
- Handles tasks such as insertion of an element into a vector, extraction of an element from a vector, vector shuffles, vector slides etc.,
  - `vdest[i] <-- vsrc[i + slide_amount]`
- **Possible area of improvement:** Vector Reduction operations are given in the latest RISC-V spec. These have not been implemented in Ara. Vector Reductions are managed using SLDU.

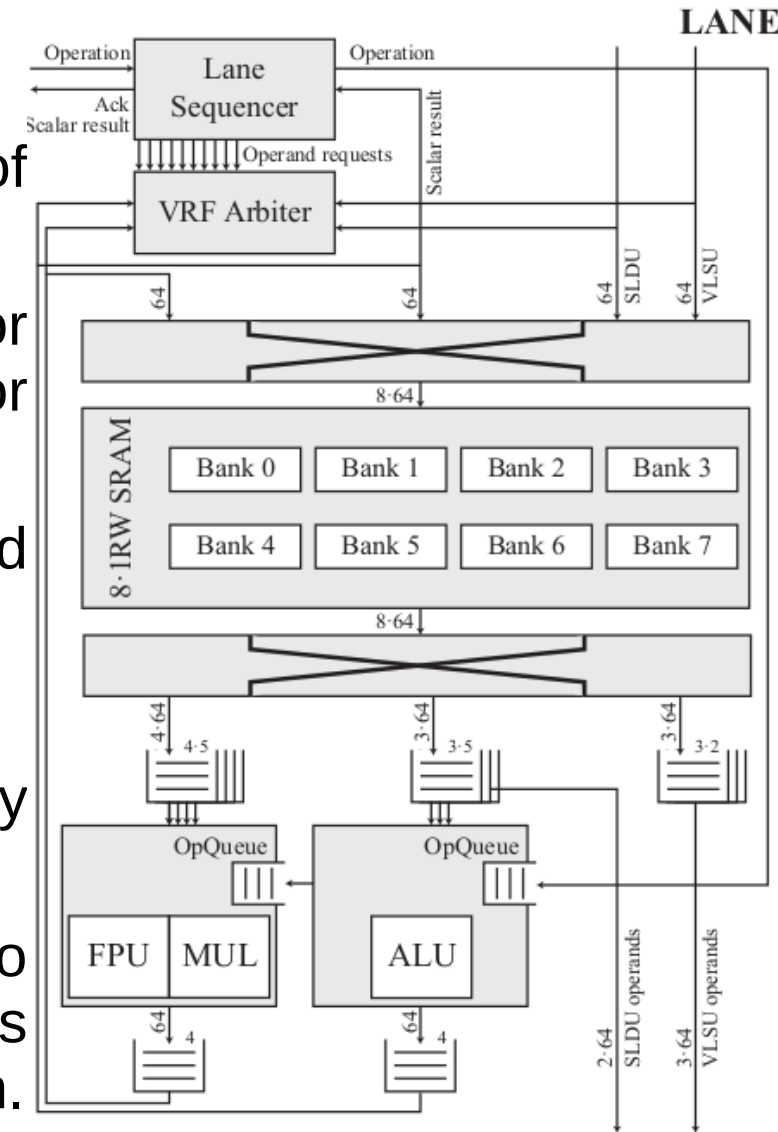
# Vector Load Store Unit (VLSU)

- Ara has a single memory port.
- VLSU has an address generator to generate address sequences for :
  - Unit Stride loads and stores, which access a contiguous chunk of memory
  - Constant Stride memory Operations, which access memory address spaced by a fixed offset.
  - Scatters and Gathers , which uses a vector of offsets to allow general access patterns.
- This block converts unit-stride memory requests into burst requests through AXI interface, instead of accessing all elements individually.



## Vector Lanes : Organization

- Ara can be configured with variable number of identical lanes
- Each lane has its own *lane-sequencer* for keeping track of upto 8 parallel vector instructions.
- Each lane contains part of Ara's whole *VRF* and *execution units*.
- VRF arbiter* controls access to VRF
- Inter-lane communication is only required by VLSU and SLDU
- Each lane has a command interface attached to the main sequencer, through which the lanes indicate they finished execution of an instruction.

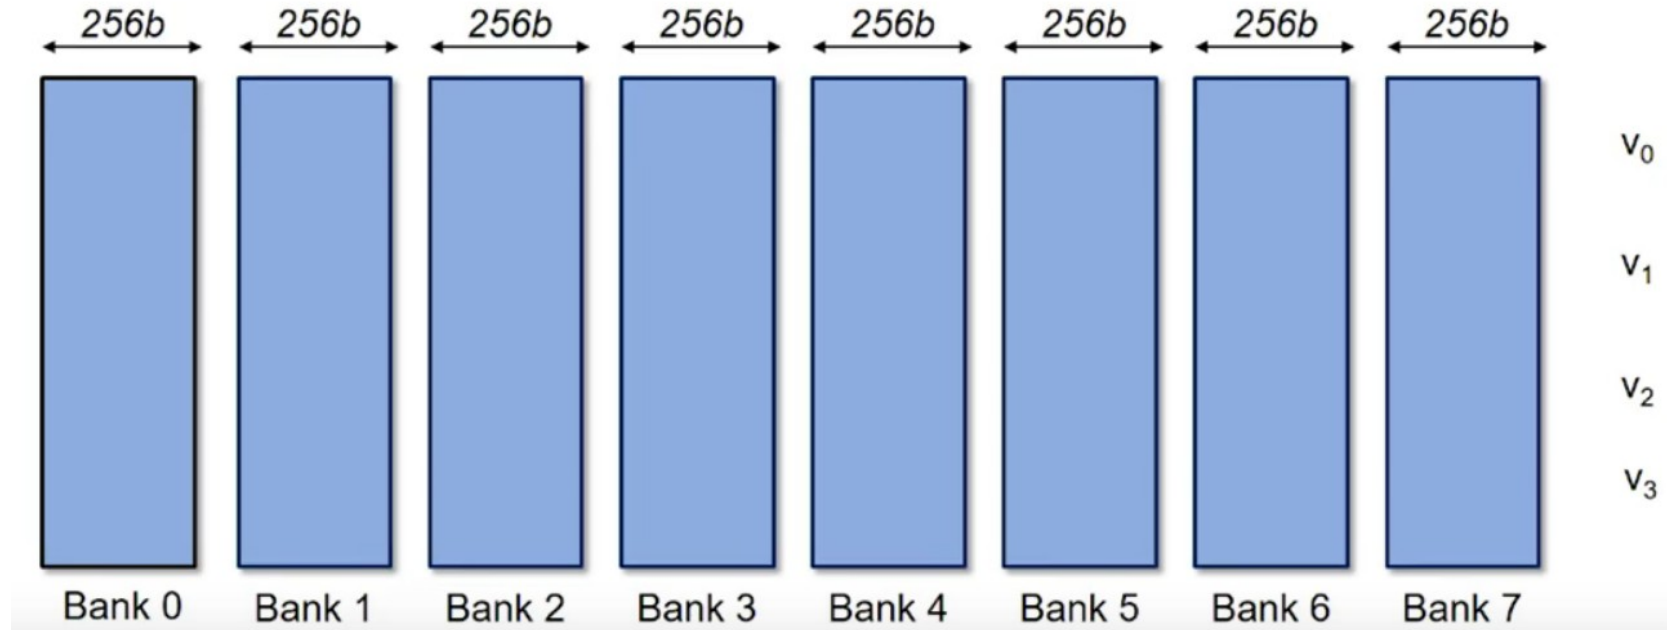


# Lane Sequencer

- Responsible for issuing vector instructions to the functional units, controlling their execution in the context of a single lane.
- Generate requests to read operands from VRF
- Upto 10 operand requests can be generated which are given to the VRF arbiter.
- In case of data hazards, the operand request rate is lowered as there is NO FORWARDING logic.

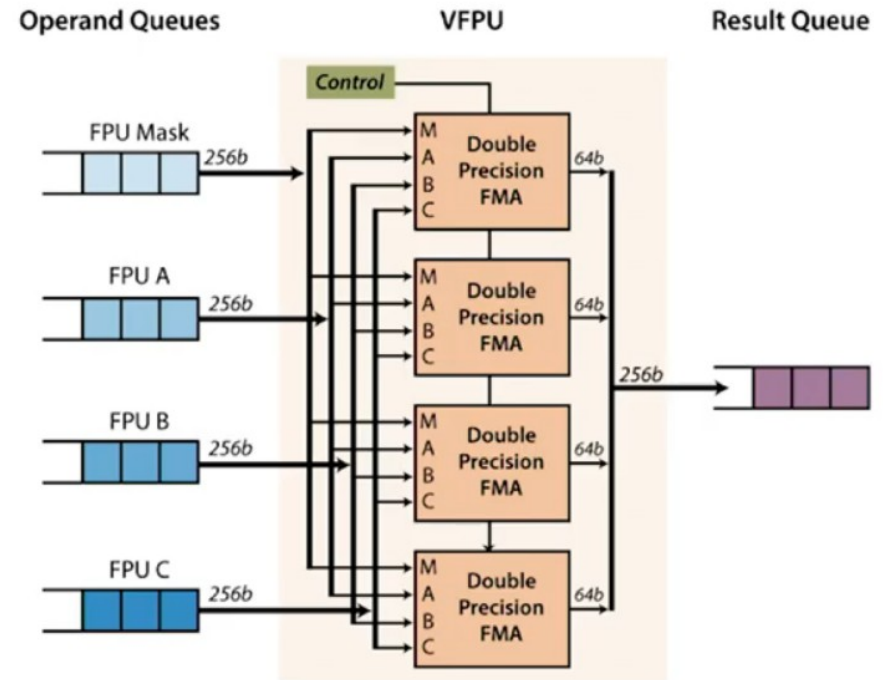
# Vector Register File

- Assume there are 4 lanes ( $N = 4$ ). Reg. File contains 256b banks. However, physically, each lane contains only 1/4 th of each bank.
- Assuming there are only 4 Vector Registers in the VRF, One bank stores 4 operands consumed in parallel by 4 lanes.



# Most Complex Instruction: Fused Multiply Add (FMA instruction)

- FMA instructions contain Four operands
- $vd[i] = \text{lsb}(\text{vmask}[i]) ? \text{vsa}[i] + \text{vsb}[i] + \text{vsc}[i] : 0;$
- FMA is pipelined (5 cycles) to meet fminconstraint
- The four lanes operate in *lockstep*
  - Low control overhead
- Each lane gets 64b operands from four 256b input FIFO buffers (A, B, C, VMASK)
  - Number of lanes determines buffer width
- 



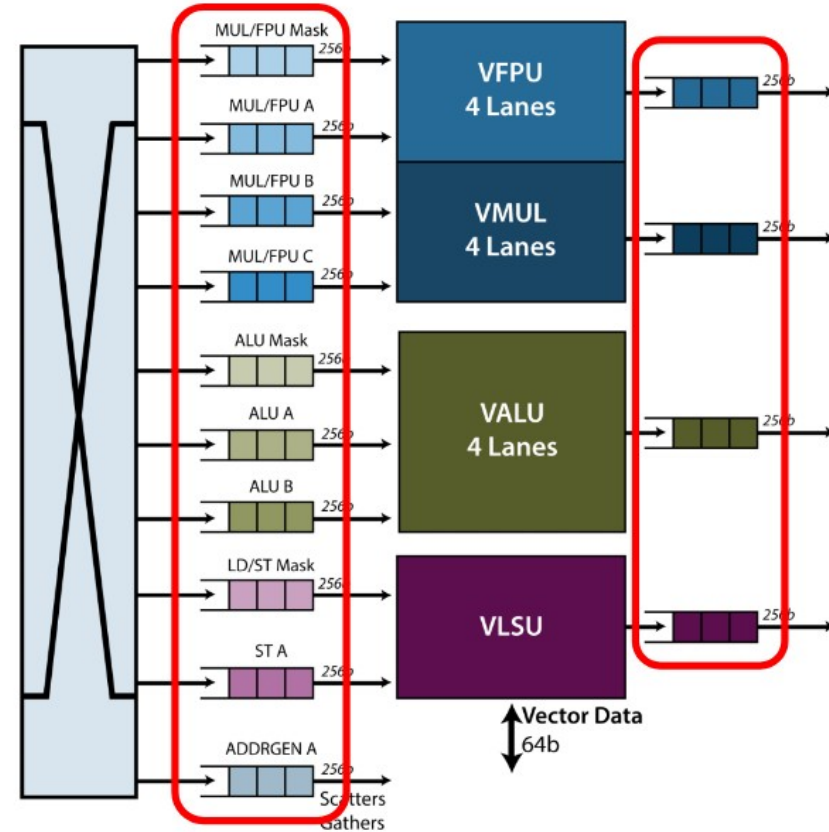


# Vector Register File

- Multi-ported Registers are essential for a VRF as several instructions in the pipeline or several data in the same instruction may access the same RF elements independently.
- Due to physical design implementation constraints, Ara's VRF is implemented as a set of single-ported (1RW) banks.
- The width of each bank is 64-bits
- In RISC-V's vector instruction, the predicated multiply and add (FMA) instruction is the worst case regarding throughput, which reads 4 operands to produce one result and takes 5 clock cycles(details in next slide)
- Therefore, in steady state, 5 banks are accessed simultaneously to sustain maximum throughput for FMA instruction.
- Ara's VRF has 8 banks per lane.


# Operand Queues

- Queues needed to sustain maximum throughput for the lock-step operation of the FUs, while hiding the latency caused by banking conflicts in the VRF
- One input queue buffer provides one operand to a the (four) lanes
- 256b (4x64b) wide entries
- One FIFO buffer per operand per multi-lane datapath unit
- 10 FIFO buffers
- Output queue buffers for output operands, one per multi-lane datapath unit

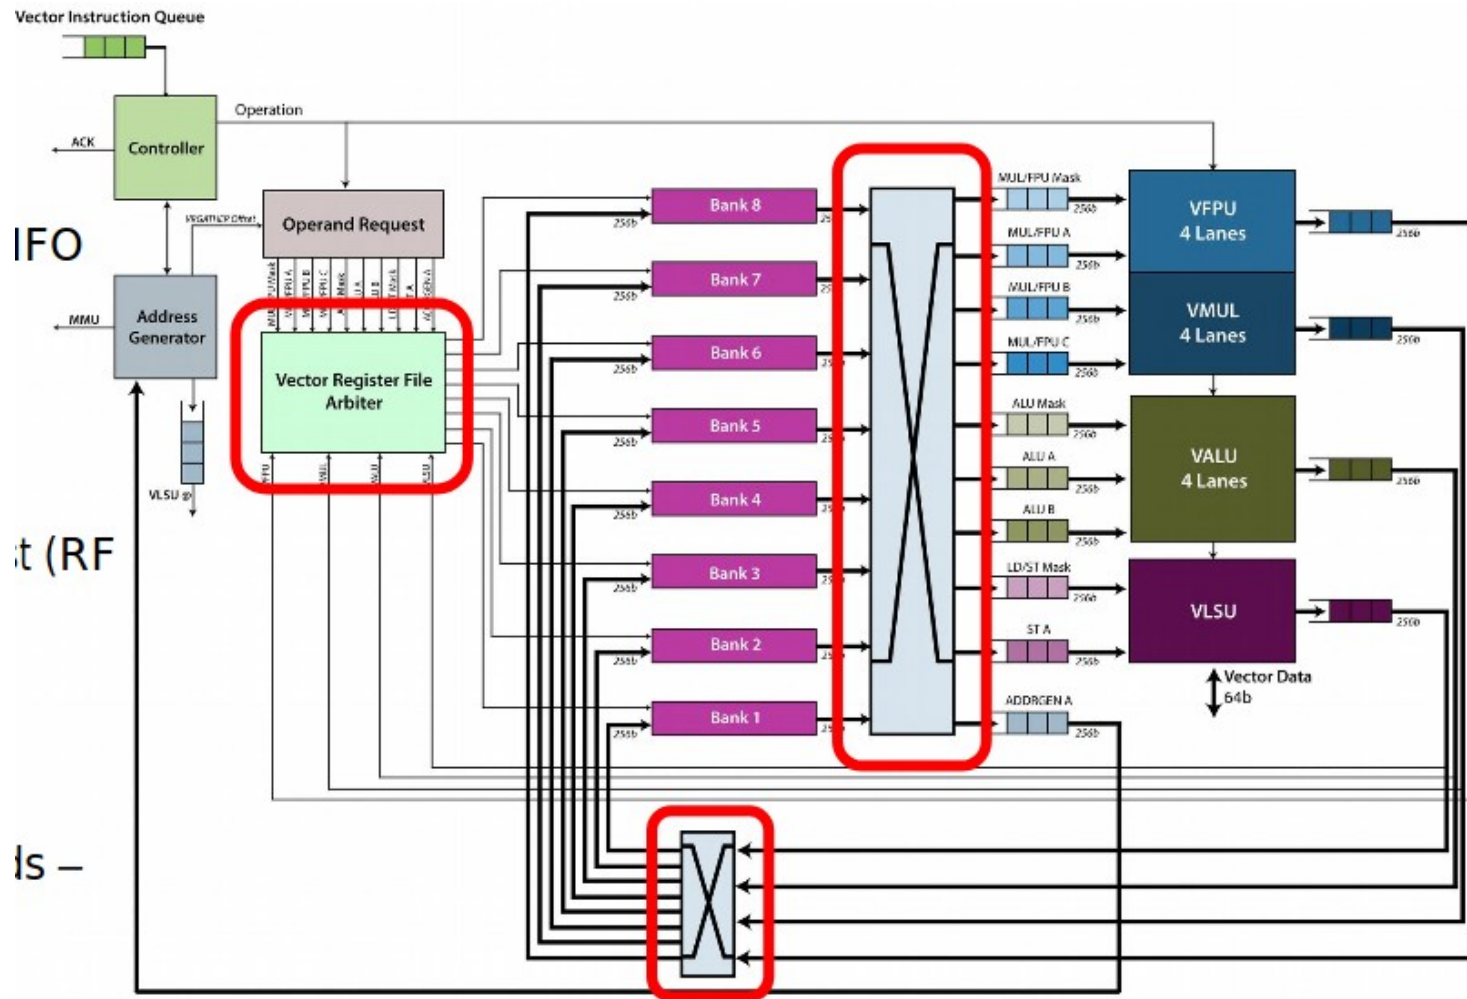


Operand Q : Courtesy : Ara RISCv Summit

# Data Organization in VRF



# Vector Register File and Operand-Deliver Interconnect



# Vector Register File and Operand-Deliver Interconnect

- All-to-all input log-interconnect
  - 256b wide (64bx4)
  - 8-source (VRF banks) x 10-dest (FIFO buffers)
  - Registered boundaries (for timing)
- All-to-all output log-interconnect
  - 256b wide (64bx4)
  - 4-source (out FIFO buffers) x 8 dest(RF banks)
- Fixed-priority arbiter
  - VRF is built as 1RW SRAM bank
  - $PM > PA > PB > PC$
  - Writes have lower priority than reads –unless output queue is full

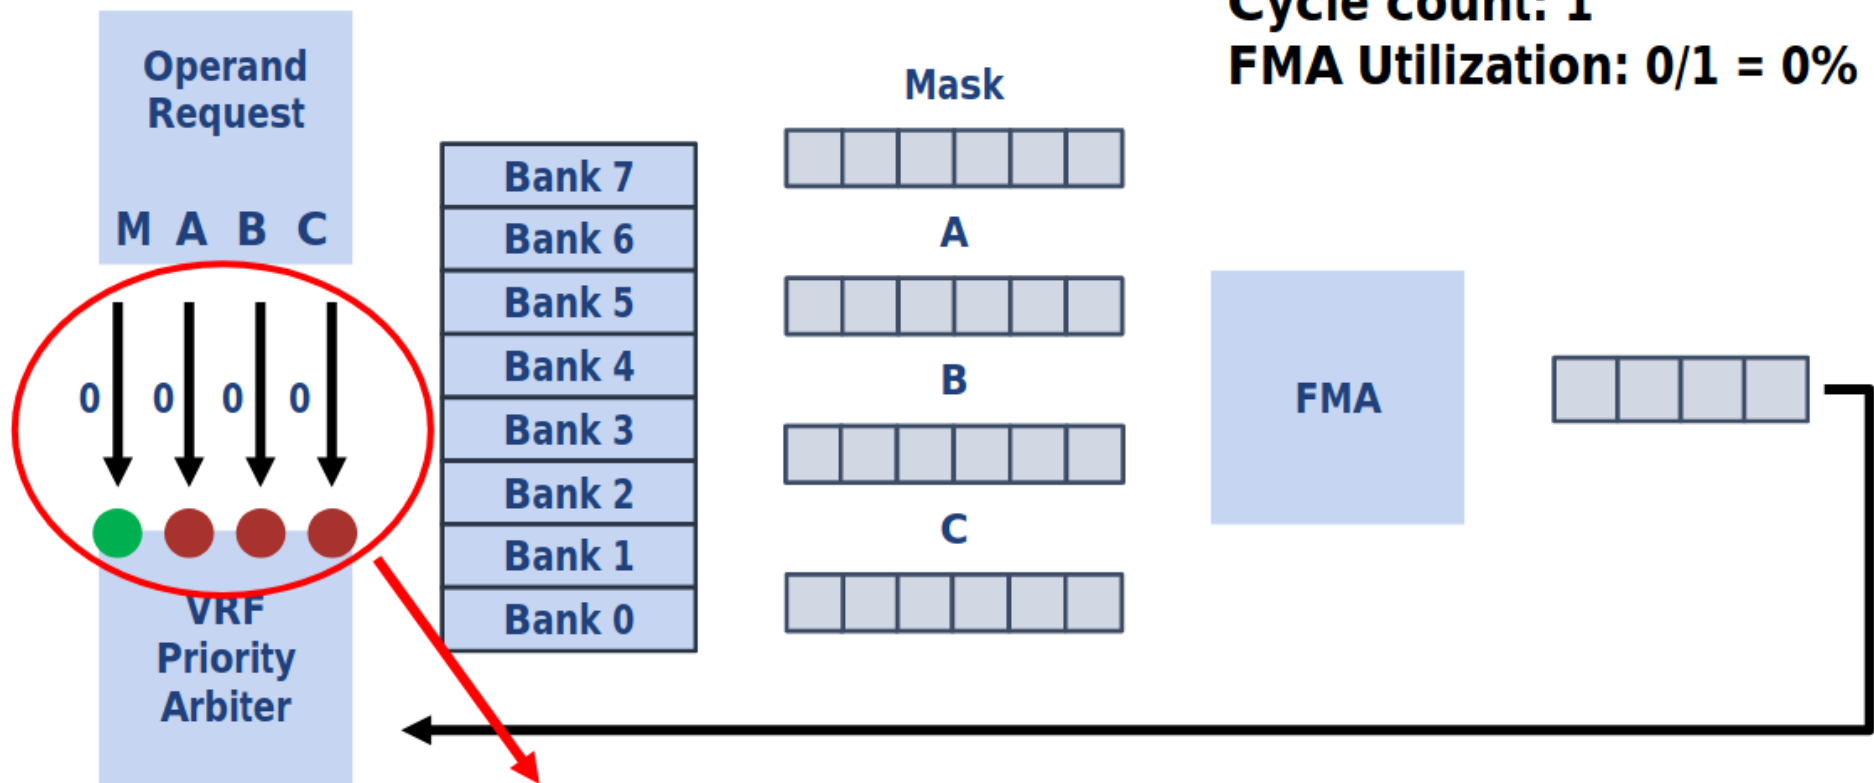
# Computation Cycle – Execution of an FMA Instruction

- Consider the execution of the following instruction:

`vmadd vd, vsa, vsb, vsc, vsmask`

- We take a vector of length = 256
- What is the execution time?
- Ideally, with 4 lanes, it should be  $256/4 = 64$
- However, due to overheads due to pipeline start, register bank clashes, it will be slightly more.

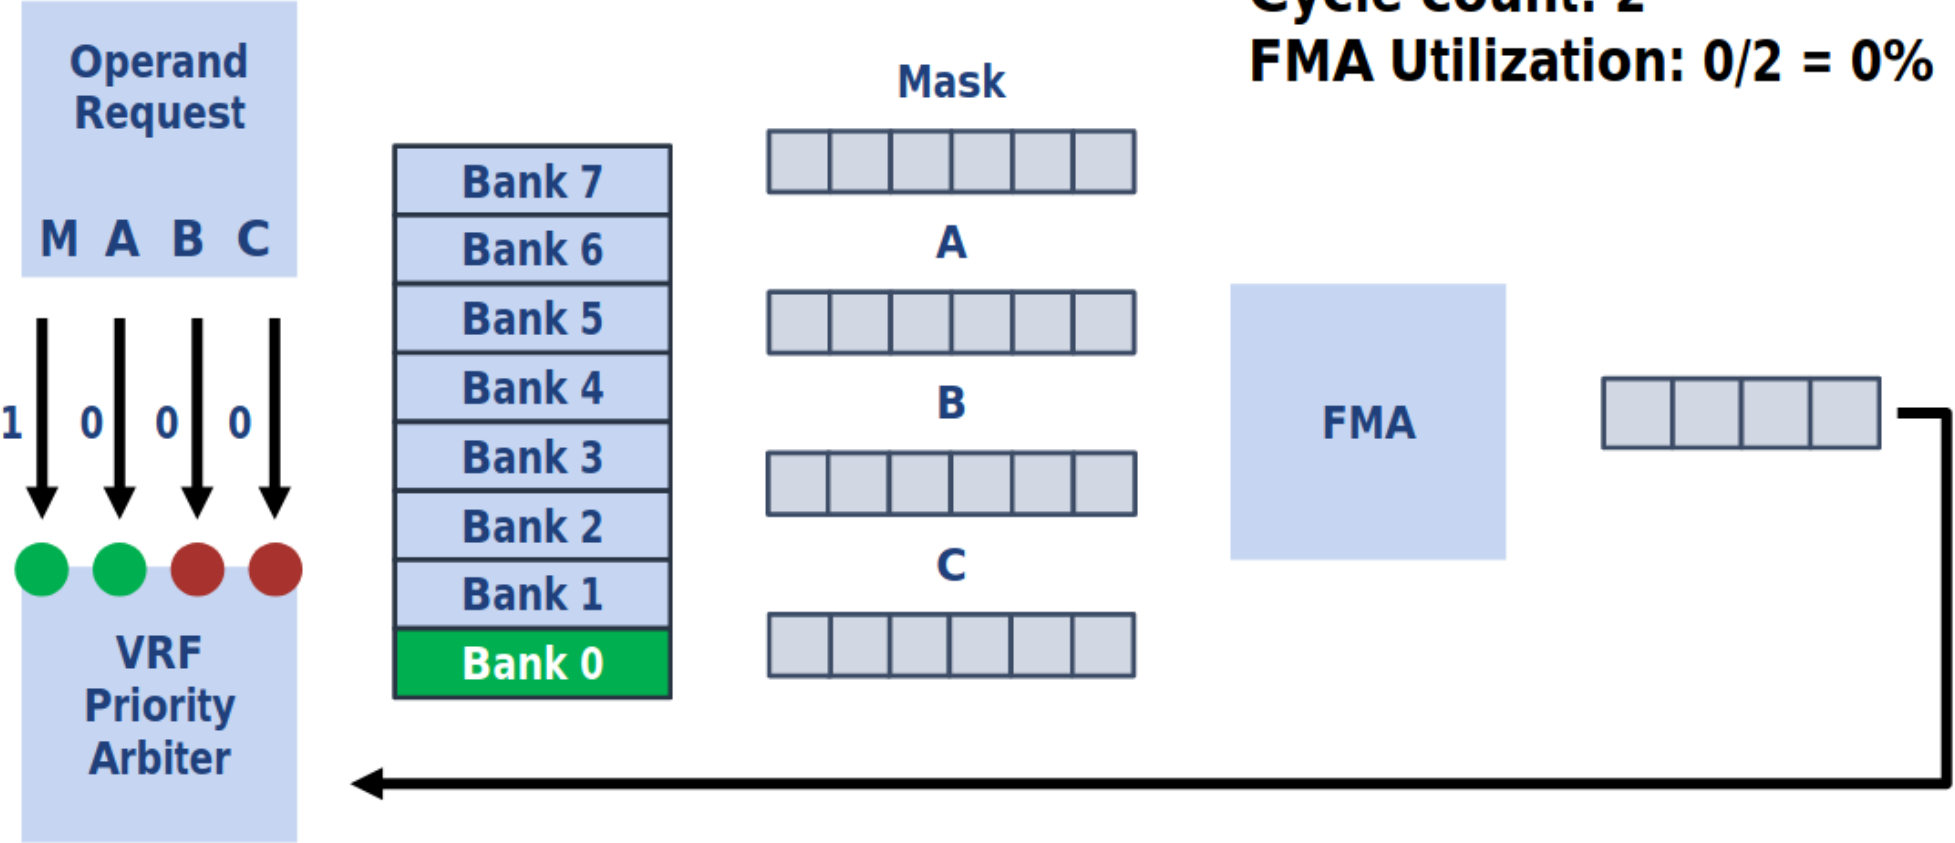
# Execution of a FMA instruction



The first 4 elements of all 4 operands are in Bank 0  
3 access stalled due to banking conflicts

# Execution of a FMA instruction

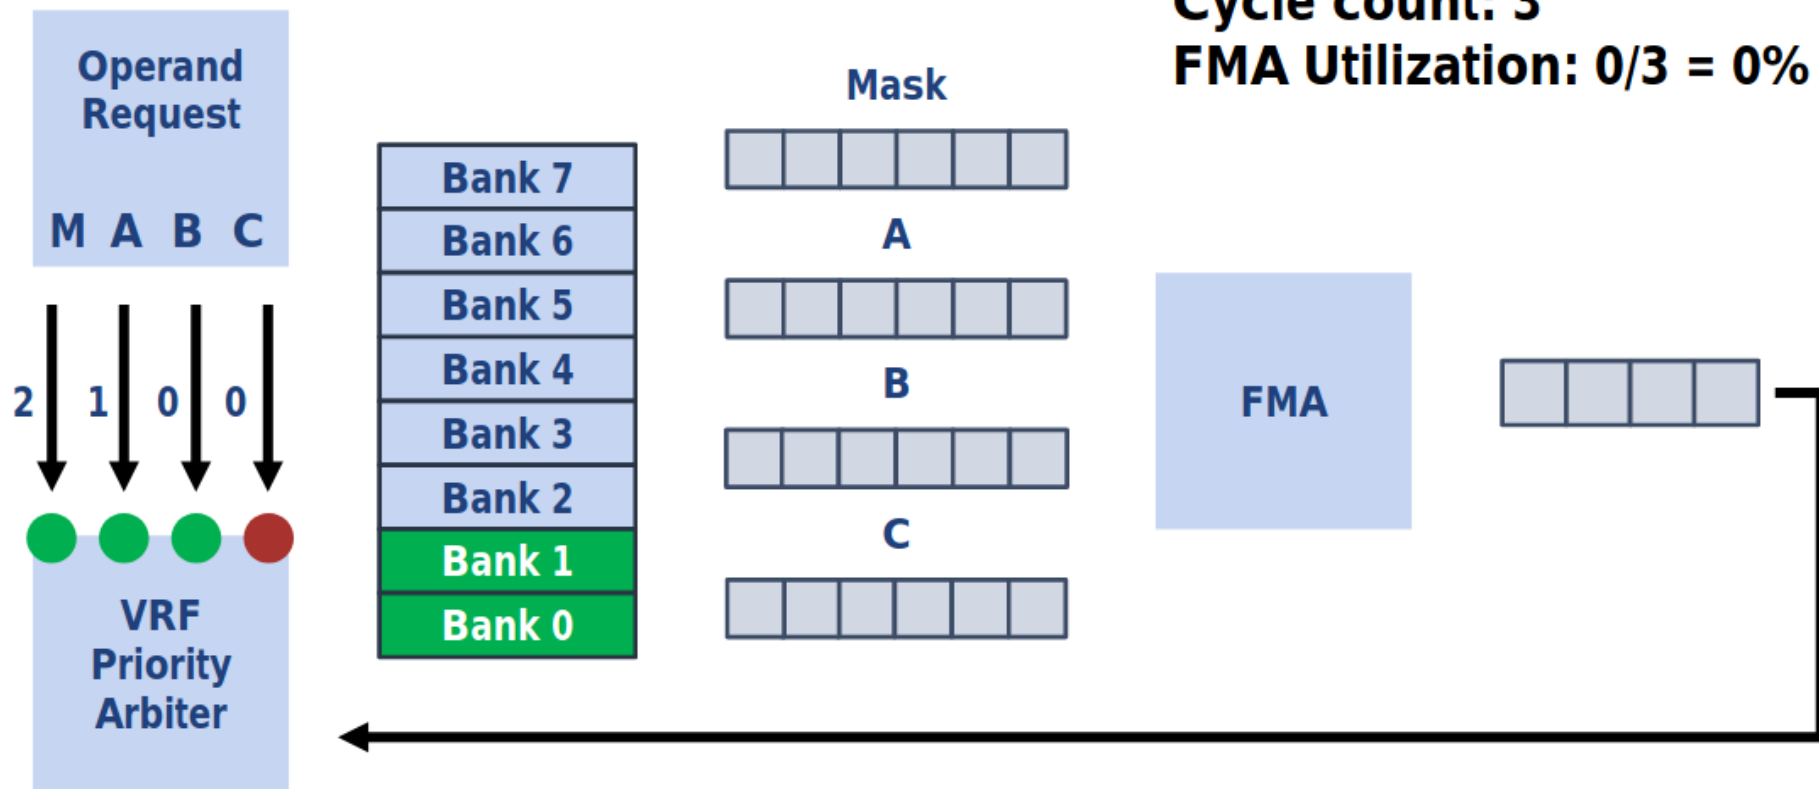
Cycle count: 2  
FMA Utilization:  $0/2 = 0\%$

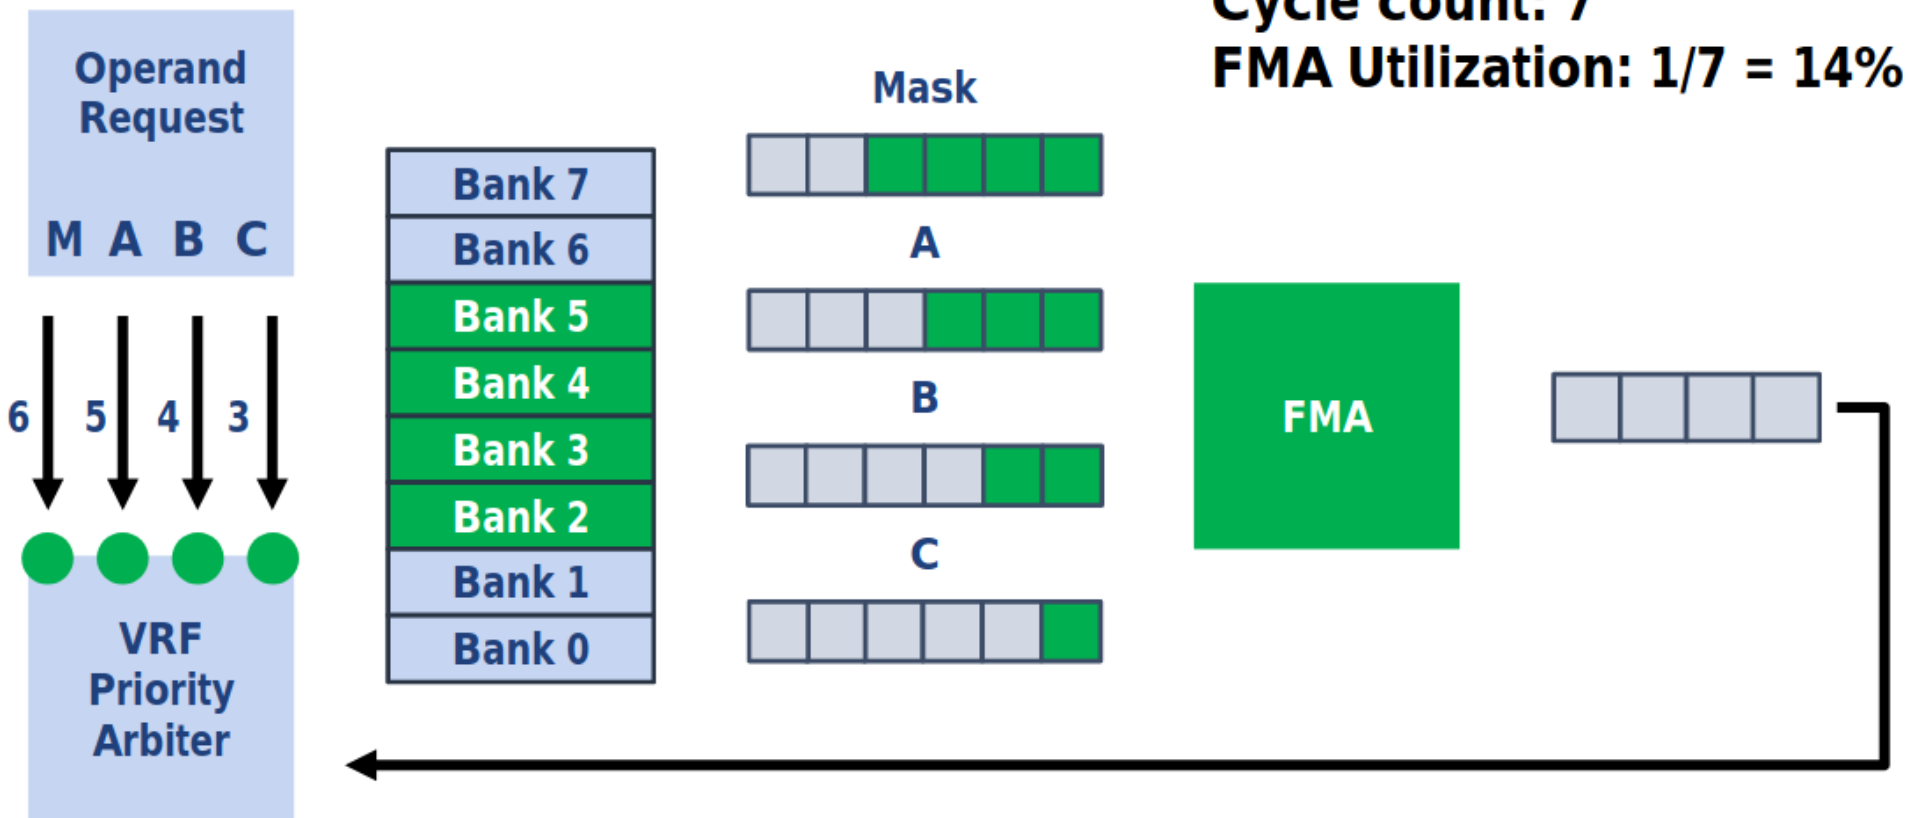


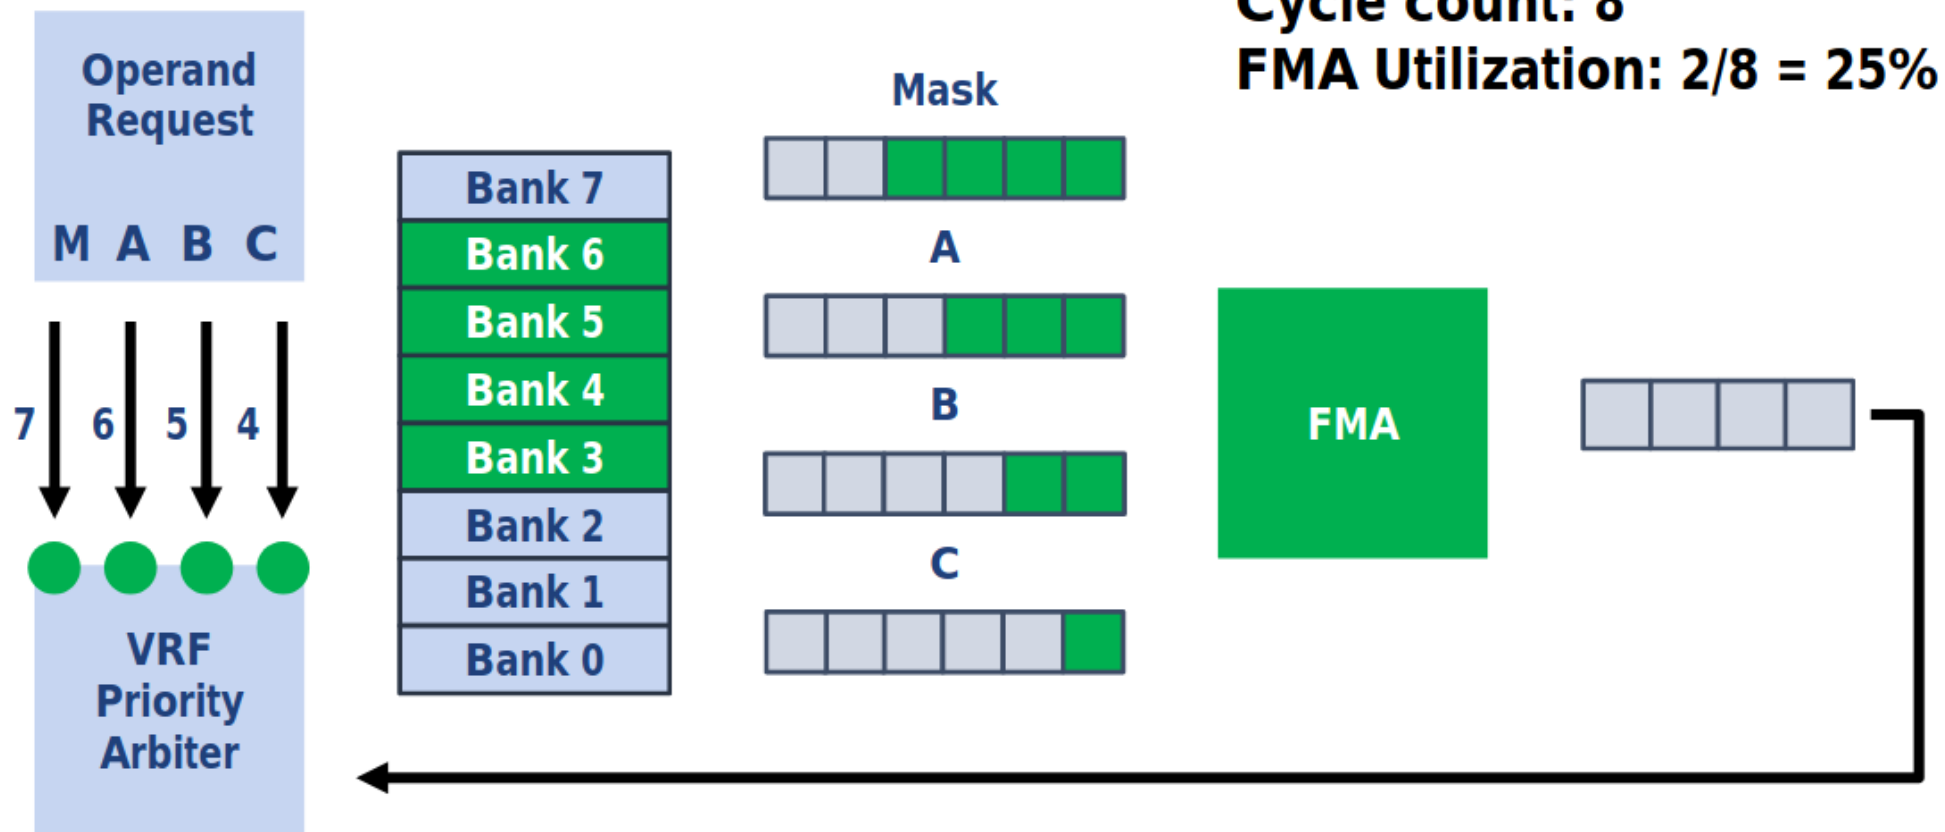


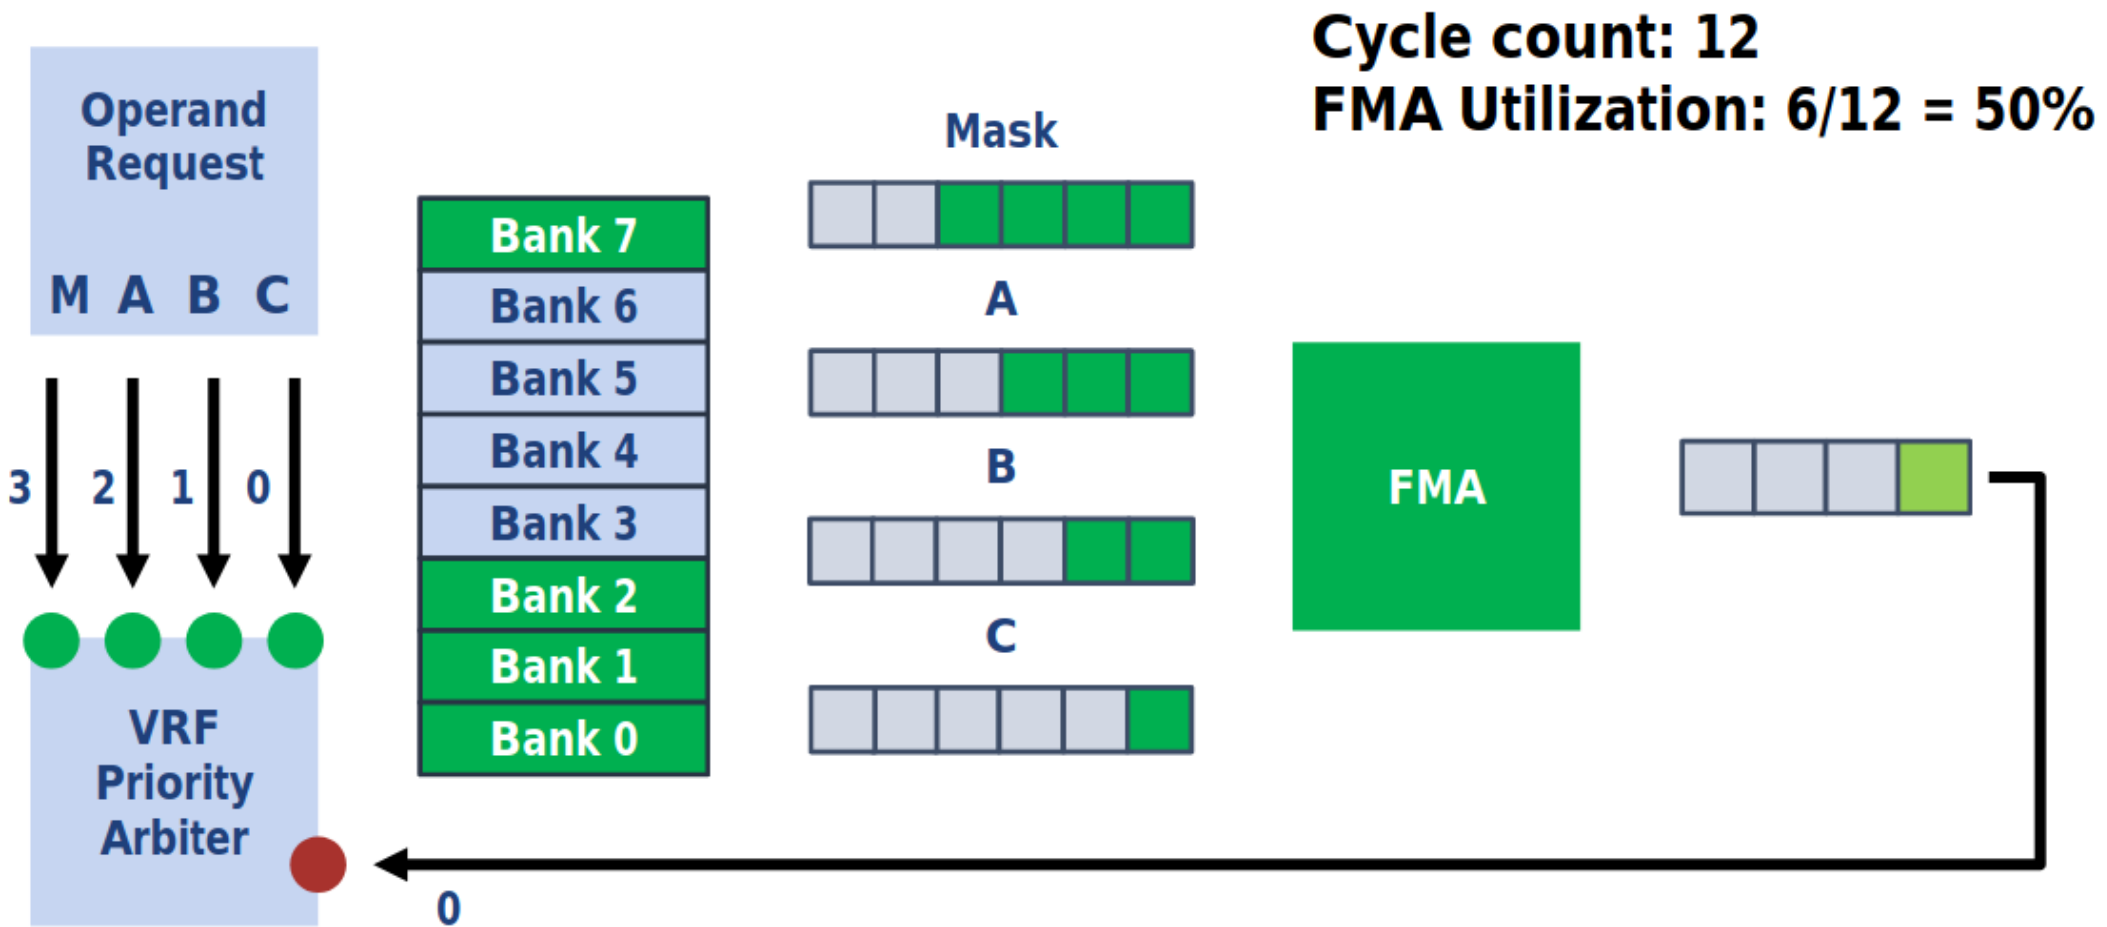
# Execution of a FMA instruction

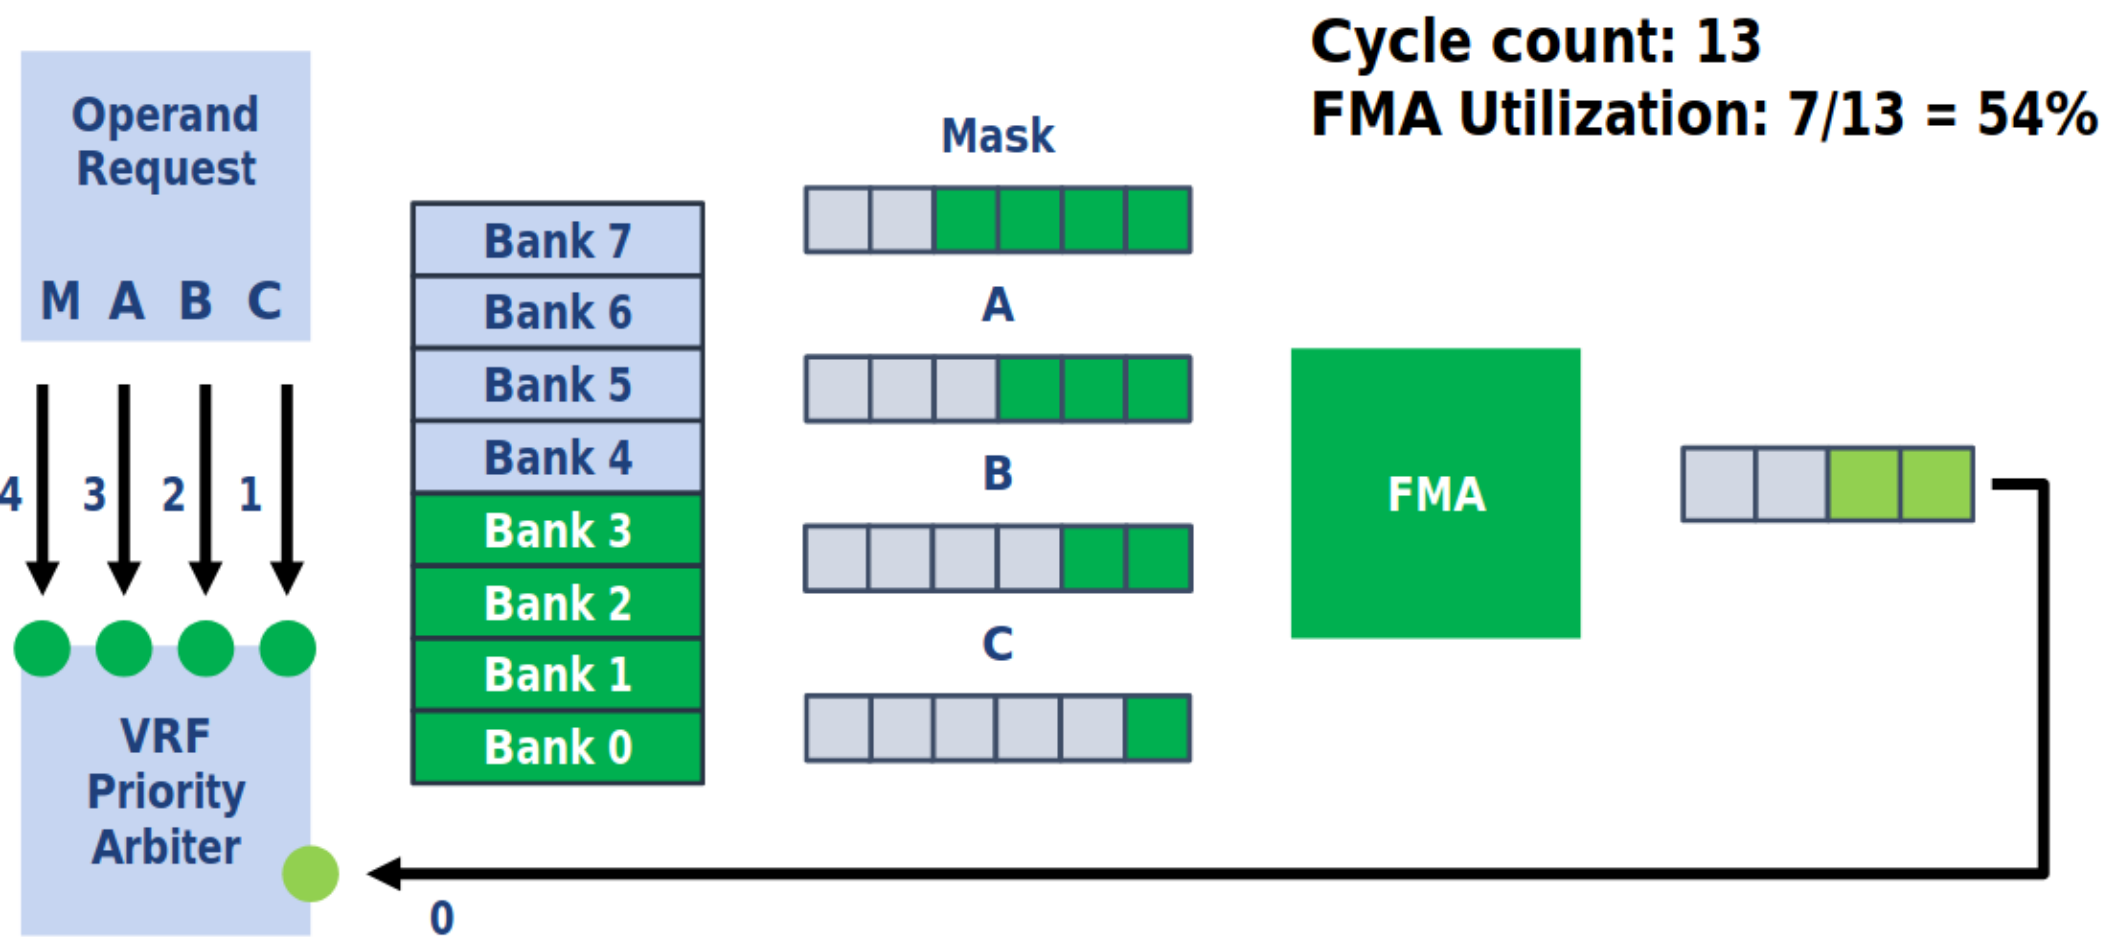
Cycle count: 3  
FMA Utilization:  $0/3 = 0\%$

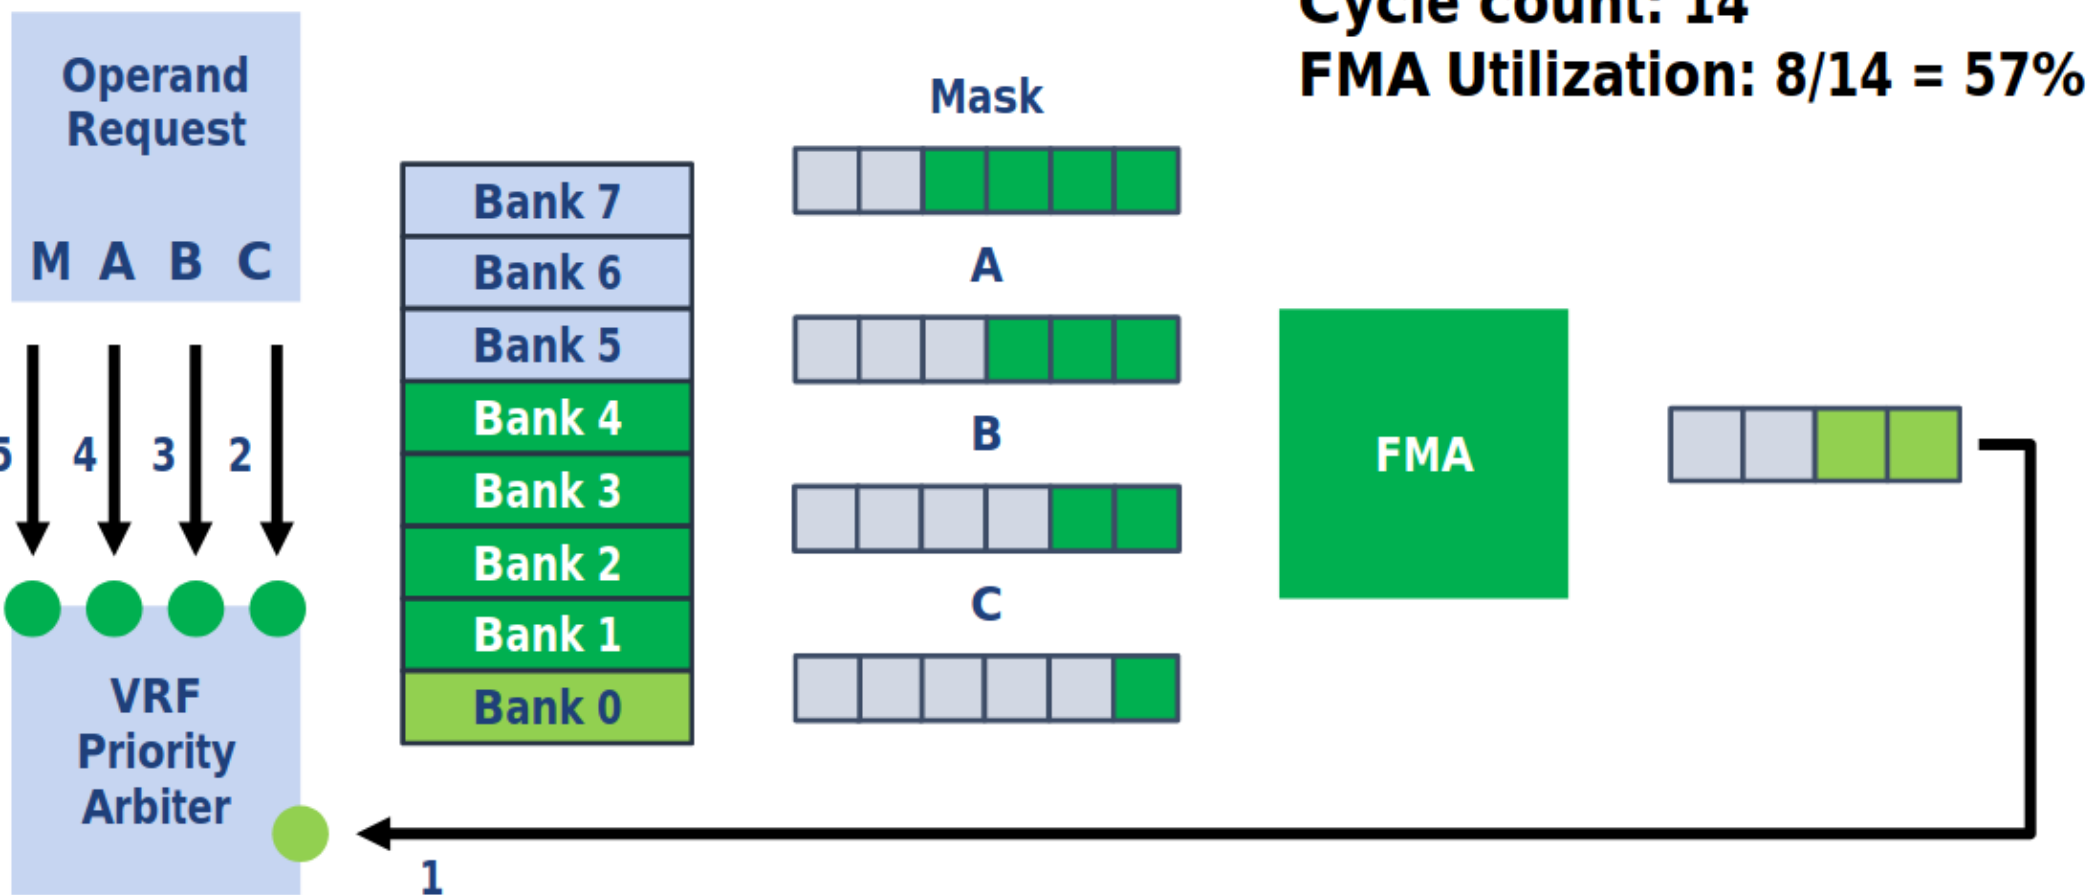


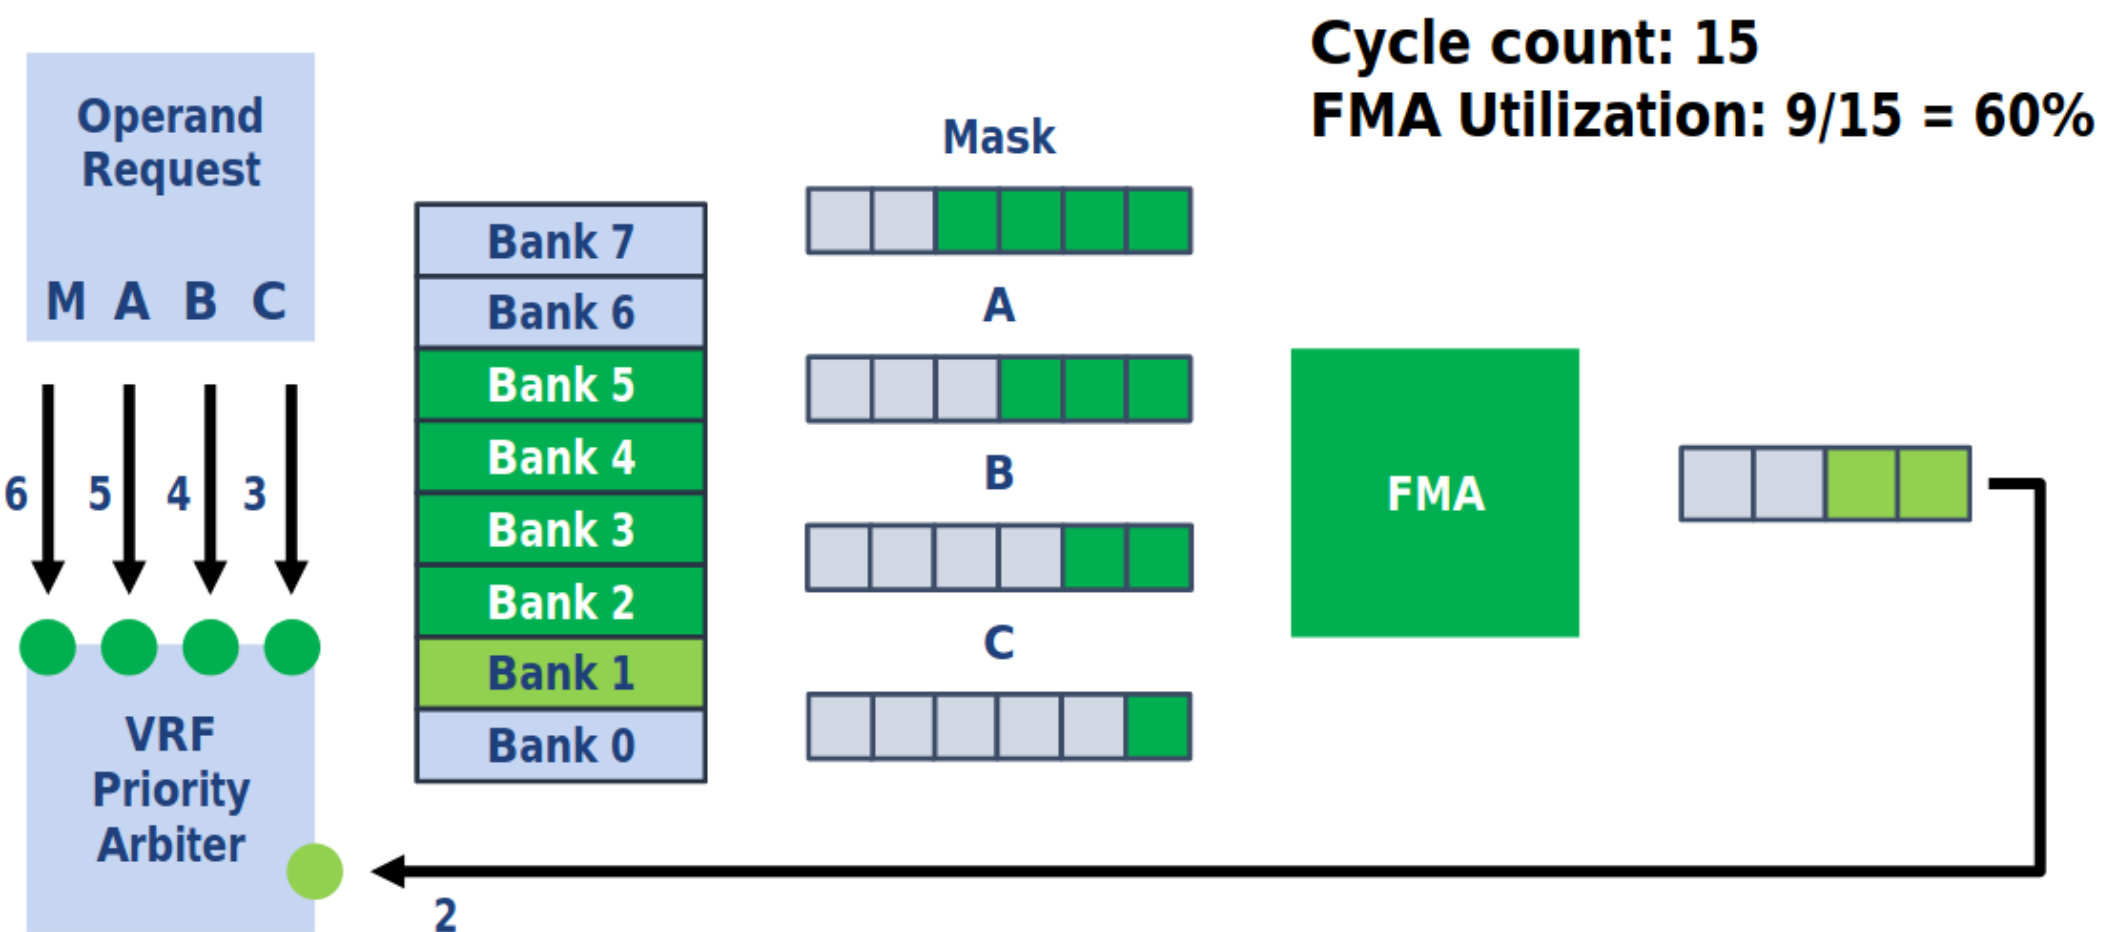












Operand  
Request

M A B C

Mask

Bank 7

Bank 6

Bank 5

Bank 4

Bank 3

Bank 2

Bank 1

Bank 0

A

B

C

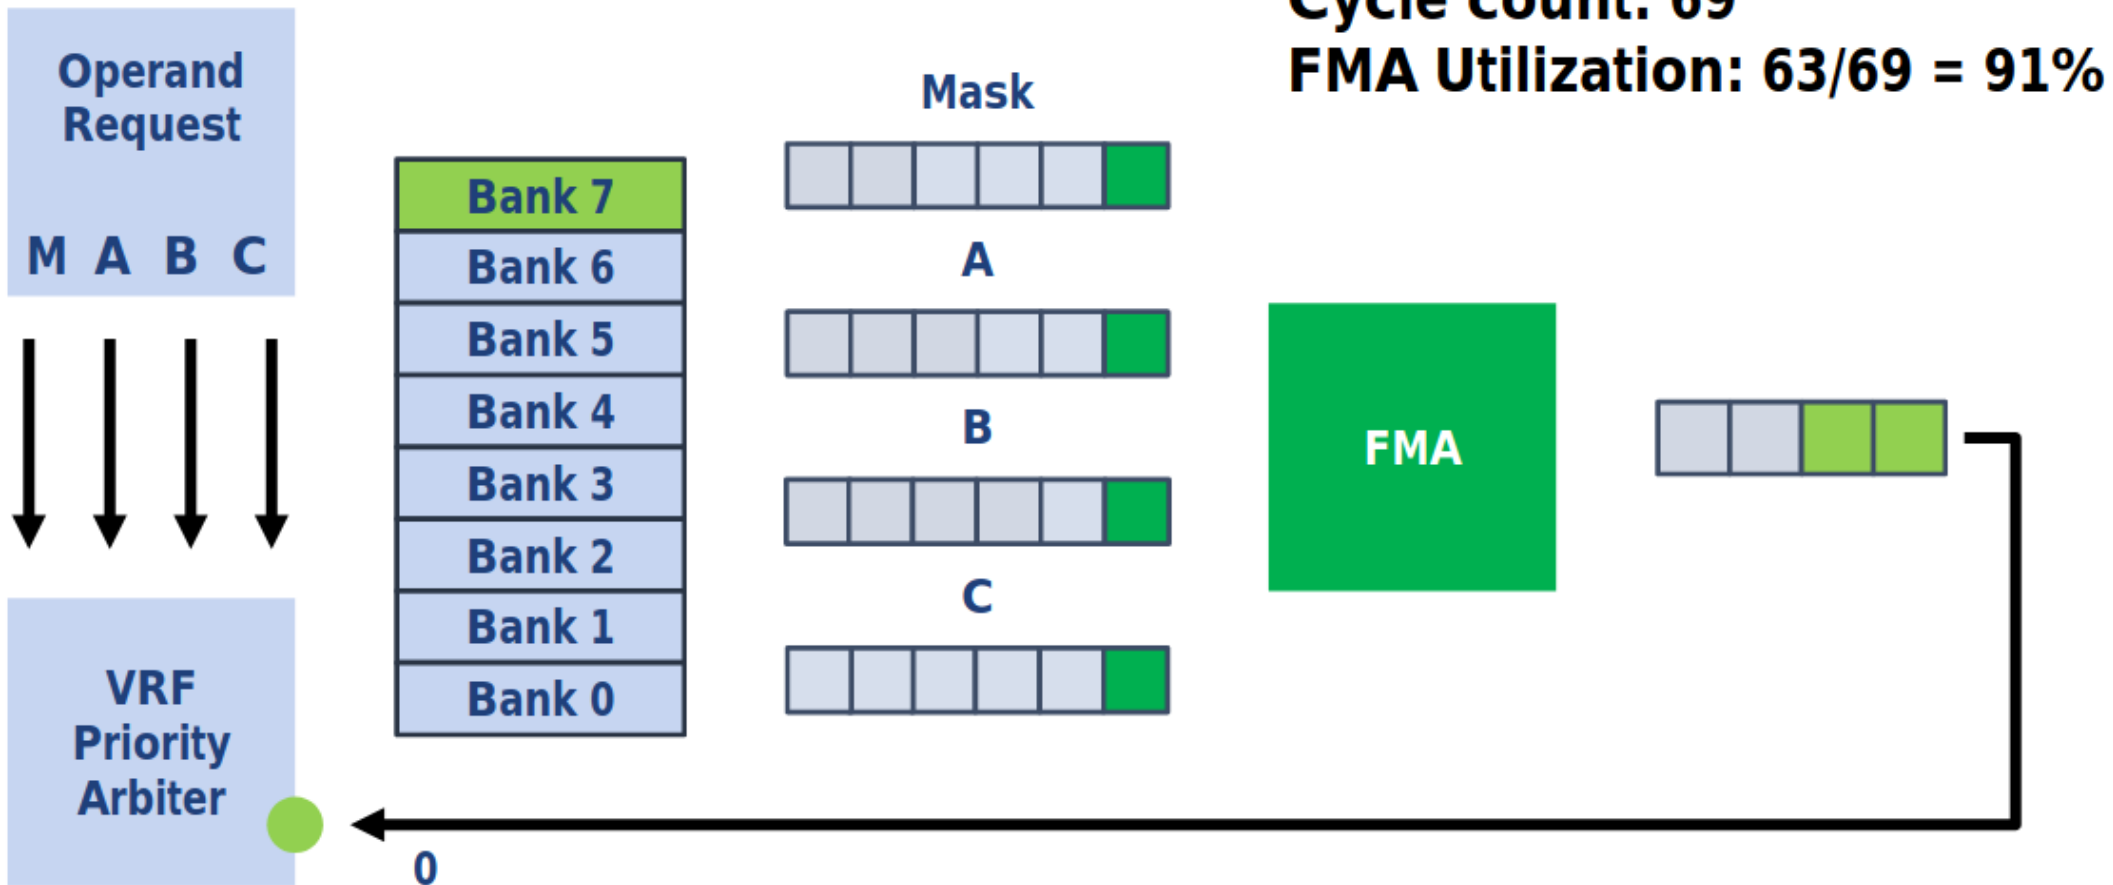
Cycle count: 69

FMA Utilization:  $63/69 = 91\%$

FMA

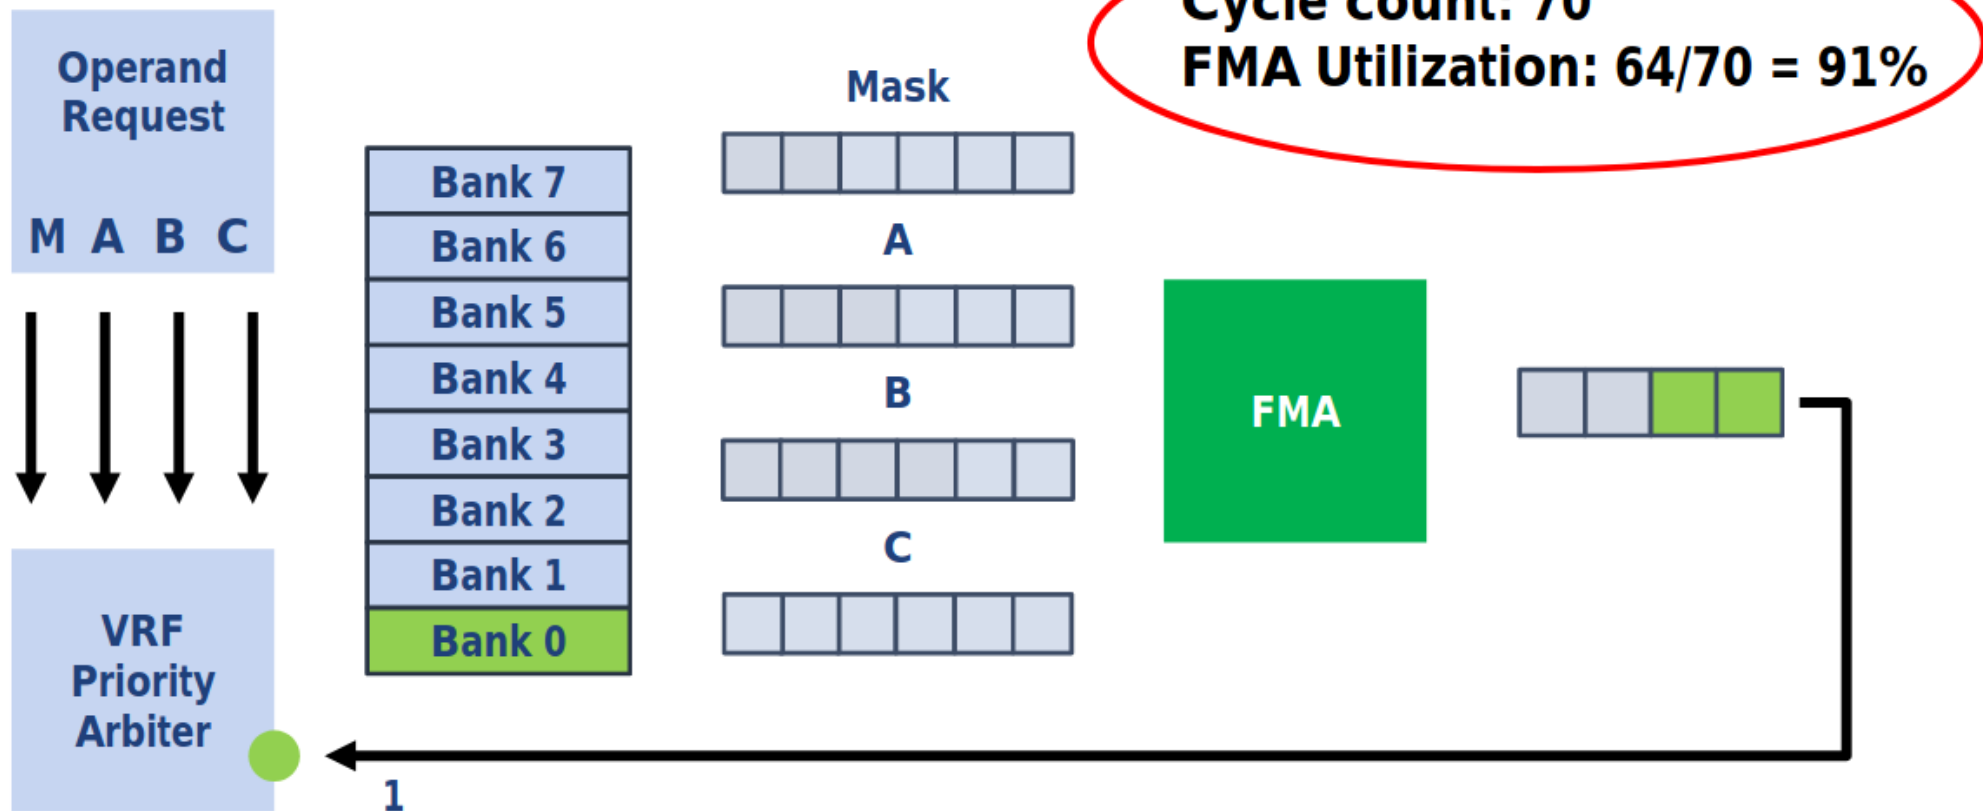
VRF  
Priority  
Arbiter

0





# Execution of a FMA instruction



# Benchmarking

The relationship between processor performance and memory bandwidth can be analyzed with the *roofline model*.

- Depicts the peak achievable performance (in OP/cycle) as a function of the arithmetic intensity  $I$ , defined as the algorithm-dependent ratio of operations per byte of memory traffic.
- According to this model, computations can be either compute-bound or memory bound
- The peak performance is achievable only if the algorithm's arithmetic intensity, in operations per byte, is higher than processor's performance per memory bandwidth ratio.
- For Ara, it enters compute-bound region when arithmetic intensity  $> 0.5$  DP-FLOP/B.




Fig. 5. Performance results for the matrix multiplication  $C \leftarrow AB + C$ , with different number of lanes  $\ell$ , for several  $n \times n$  problem sizes. The bold red line depicts a performance boundary due to the instruction issue rate. The numbers between brackets indicate the performance loss, with respect to the theoretically achievable peak performance.

# Three Benchmarks

## • Matrix Multiplication

- $C \leftarrow AB + C$  where all matrices are  $n \times n$  having DP elements
- The algorithm requires  $2n^3$  floating-point operations (one FMA is considered as 2 operations) and  $32n^2$  Bytes of memory transfers.
- Therefore, the algorithm has **arithmetic intensity  $2n^3 / 32n^2 = n/16$  DP-FLOPS/B**

## • Convolution(DCONV)

- Example from GoogLeNet CNN, Input Image =  $112 \times 112 \times 3$ ; kernel =  $7 \times 7 \times 3 \times 64$
- This results in a total of  $7 \times 7 \times 3 \times 64 \times 112 \times 112$  FMAs, or 236 DP-MFLOPS
- Memory :  $(3 \times 118 \times 118 + 7 \times 7 \times 3 \times 64 + 64 \times 112 \times 112) \times 8$  Bytes = 6.44 MiB
- Therefore, **arithmetic intensity = 34.9 DP-FLOPS/B =>Heavily Compute Bound**

## • DAXPY

- $Y \leftarrow aX + Y$  where  $X, Y$  are  $n$ -dim vectors and  $a$  is a DP scalar
- Requires  $2^n$  operations ( $n$  FMAs) and  $24n$  bytes of memory transfers
- Therefore, this algorithm has **arithmetic intensity =  $1/12$  DP-FLOPS/B =>Memory Bound**

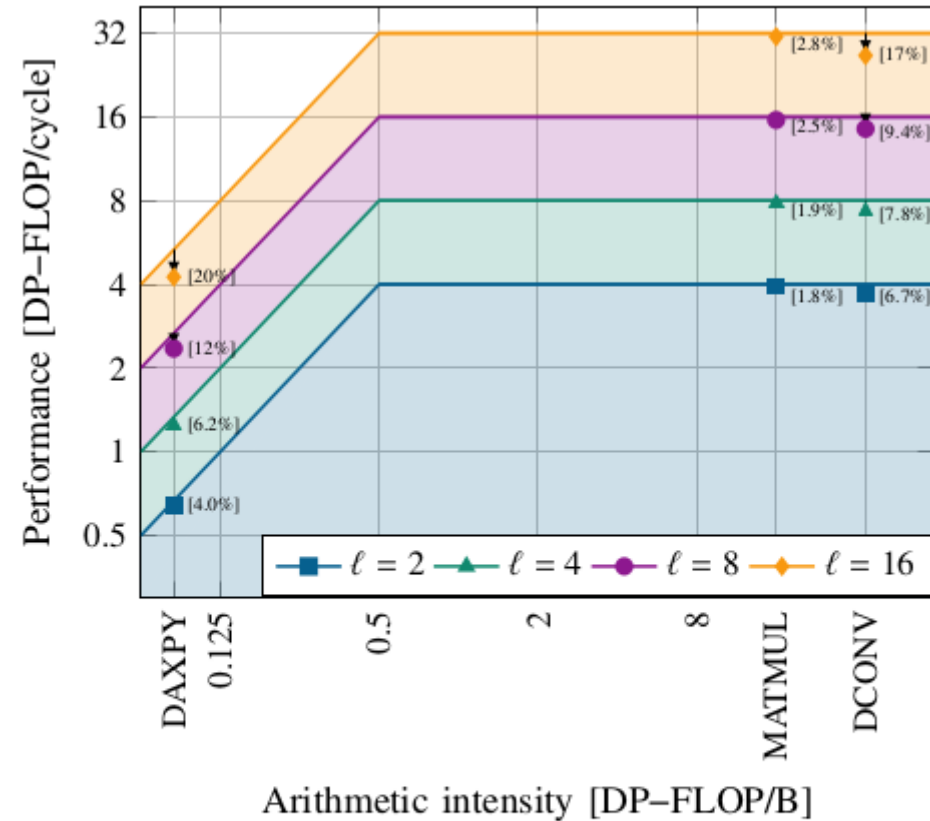


Fig. 6. Performance results for the three considered benchmarks, with different number of lanes  $\ell$ . AXPY uses vectors of length 256, the MATMUL is between matrices of size  $256 \times 256$ , and CONV uses GoogLeNet's sizes. The numbers between brackets indicate the performance loss, with respect to the theoretically achievable peak performance.

# Week 11: Oct 9<sup>th</sup> 2020

## **Actions**

1. Study RISC-V Vector Extension

## **Work Done**

1. Study RISC-V Vector Extension Spec (7 out 20 Chapters)

## Goals for Standard RISC-V “V” Extension

- Efficient and Scalable to all reasonable design-points
  - Low Cost Micro-controller or high-performance supercomputer
  - In-order, decoupled, or out-of-order microarchitectures
  - Integer, fixed-point, or floating-point data types
- Compiler Support for Auto-vectorization
- All instructions to fit into Standard 32-bit encoding space
- Be base for future vector++ extensions such as Crypto, Machine Learning extensions on top of RISC-V “V” extension.

**Source:** *RISCV “V” Extension Proposal by Krste Asanovic, UC-Berkeley and SIFive, 2016*

# V Key features

- Cray-style vectors
  - “The right way” to exploit SIMD parallelism
  - Packed-SIMD, GPU etc., are inefficient in terms of number of instructions in the instruction set, number of instructions executed and lack of portability. (Source: Computer Organization RISC-V Edition , Chapter 8)


<i>ISA</i>	<i>MIPS-32 MSA</i>	<i>IA-32 AVX2</i>	<i>RV32V</i>
Instructions (static)	22	29	13
Instructions per Main Loop	7	6	10
Bookkeeping Instructions	15	23	3
Results per Main Loop	2	4	64
Instructions (dynamic n=1000)	3511	1517	163

*Source: Comp Org. Ch 8*

- Implementation-dependent vector length
  - Same binary runs on different CPUs with different hardware vector lengths
- Reconfigurable Vector Register File (allow different vector lengths and element sizes)
- Mixed Precision Support (can operate with operands with different widths in the same op)

# V Extension State


- Standard RISC-V scalar x and f registers



- Vector Configuration
  - Vector Length Register **vl**
  - Vector type register **vtype**
  - Other Control Registers
    - vstart** (For trap handling)
    - vrmsat** - Fixed point rounding mode/ saturation

- Upto 32 vector registers v0-v31 of atleast 4 elements each, with variable bits/element(size in powers of 2)

Vector data registers



- One Predicate Register - **v0** for containing vector masks

# Vector Unit Implementation

## Dependent Parameters

- **ELEN** : Size of largest element in bits
- **VLEN** : No. Of bits in each vector register
  - $VLEN \geq ELEN$
- **SLEN** : Striping Distance in bits
  - Sets how many bits are packed contiguously into one vector register before moving to the next vector register in the *group*.
  - The striping length SLEN for an implementation is set to optimize the tradeoff between datapath wiring for mixed-width operations and buffering needed to corner-turn wide vector unit-stride memory accesses into parallel accesses for the vector register file.
  - This is generally the width of the memory port.
  - $VLEN \geq SLEN \geq ELEN$




# Some Microarchitecture Design Points

Name	Issue Policy	Issue Width	VLEN (bits)	Datapath (bits)	VLEN/Datapath (beats)
Smallest	InO	1	32	32	1
Simple	InO	1	512	128	4
InO-Spatial	InO	2	128	128	1
OoO-Spatial	OoO	2-3	128	128	1
OoO-Temporal	OoO	2-3	512	128	4
OoO-Server	OoO	3-6	2048	512	4
OoO-HPC	OoO	3-6	16384	2048	8

**Source:** Vector Extension 0.7 Slides , Krste Asanovic, UC Berkeley, RISC-V Summit - 2019

# Vector Type Register



Bits	Contents
1:0	vlmul[1:0]
4:2	vsew[2:0]
6:5	vediv[1:0]
XLEN-1:7	Reserved (write 0)

vsew[2:0]			SEW
0	0	0	8
0	0	1	16
0	1	0	32
0	1	1	64
1	0	0	128
1	0	1	256
1	1	0	512
1	1	1	1024

**vsew[2:0]** field encodes standard element width (SEW) in bits of elements in vector register ( $SEW = 8 * 2^{vsew}$ )

**vlmul[1:0]** encodes vector register length multiplier ( $LMUL = 2^{vlmul} = 1-8$ ) *(v0.9 adds "fractional LMUL" < 1)*

**vediv[1:0]** encodes how vector elements are divided into equal sub-elements ( $EDIV = 2^{vediv} = 1-8$ )

## Example Vector Register Layouts

VLEN=32b

3	2	1	0	SEW
3	2	1	0	8b
1		0		16b
			0	32b

VLEN=64b

7	6	5	4	3	2	1	0	SEW
7	6	5	4	3	2	1	0	8b
3		2		1		0		16b
1				0				32b
						0		64b

VLEN=128b

F	E	D	C	B	A	9	8	7	6	5	4	3	2	1	0	SEW
F	E	D	C	B	A	9	8	7	6	5	4	3	2	1	0	8b
7		6		5		4		3		2		1		0		16b
3				2				1				0				32b
1								0								64b
0															128b	

VLEN = 256b

1F	1E	1D	1C	1B	1A	19	18	17	16	15	14	13	12	11	10	F	E	D	C	B	A	9	8	7	6	5	4	3	2	1	0	SEW
1F	1E	1D	1C	1B	1A	19	18	17	16	15	14	13	12	11	10	F	E	D	C	B	A	9	8	7	6	5	4	3	2	1	0	8b
F		E		D		C		B		A		9		8		7		6		5		4		3		2		1		0		16b
7				6				5				4				3				2				1				0				32b
3								2								1								0								64b
1																0																128b
0																																256b

# Vector Register Grouping (lmul)

Example 1: VLEN=32b, SEW=16b, LMUL=2

```
Byte      3 2 1 0
v2*n      1  0
v2*n+1    3  2
```

Example 2: VLEN=64b, SEW=32b, LMUL=2

```
Byte      7 6 5 4 3 2 1 0
v2*n      1  0
v2*n+1    3  2
```

Example 3: VLEN=128b, SEW=32b, LMUL=2

```
Byte      F E D C B A 9 8 7 6 5 4 3 2 1 0
v2*n      3  2  1  0
v2*n+1    7  6  5  4
```

Example 4: VLEN=256b, SEW=32b, LMUL=2

```
Byte      1F1E1D1C1B1A19181716151413121110 F E D C B A 9 8 7 6 5 4 3 2 1 0
v2*n      B  A  9  8  3  2  1  0
v2*n+1    F  E  D  C  7  6  5  4
```

vlmul		LMUL	#groups	VLMAX	Grouped registers
0	0	1	32	VLEN/SEW	vn (single register in group)
0	1	2	16	2*VLEN/SEW	vn, vn+1
1	0	4	8	4*VLEN/SEW	vn, ..., vn+3
1	1	8	4	8*VLEN/SEW	vn, ..., vn+7

- Multiple registers can be grouped together so that a single instruction can operate on multiple vector registers
- Term “*Vector Register Group*” refers to one or more registers that form a single operand
- Used for :
  - Accomodate mixed-width precision ops
  - To increase efficiency by using longer vectors when fewer separate registers needed

## LMUL with shorter striping length (SLEN < VLEN)

Example 3: VLEN=128b, SLEN=64b, SEW=32b, LMUL=4

Byte	F	E	D	C	B	A	9	8	7	6	5	4	3	2	1	0	
v4*n				9				8				1				0	32b elements
v4*n+1				B				A				3				2	
v4*n+2				D				C				5				4	
v4*n+3				F				E				7				6	

Example 4: VLEN=128b, SLEN=128b, SEW=32b, LMUL=4

Byte	F	E	D	C	B	A	9	8	7	6	5	4	3	2	1	0	
v4*n				3				2				1				0	32b elements
v4*n+1				7				6				5				4	
v4*n+2				B				A				9				8	
v4*n+3				F				E				D				C	

Example 5: VLEN=256b, SLEN=128b, SEW=32b, LMUL=4

Byte	1F	E1D1C1B1A19181716151413121110	F	E	D	C	B	A	9	8	7	6	5	4	3	2	1	0
v4*n		13		12		11		10		3		2		1		0		
v4*n+1		17		16		15		14		7		6		5		4		
v4*n+2		1B		1A		19		18		B		A		9		8		
v4*n+3		1F		1E		1D		1C		F		E		D		C		

Example 6: VLEN=256b, SLEN=128b, SEW=256b, LMUL=4

[illegible]

# Setting vector configuration, vsetvli/vsetvl


The **vsetvl{i}** configuration instructions set the **vttype** register, and also set the **vl** register, returning the **vl** value in a scalar register

**vsetvli rd, rs1, e8 # Set SEW=8, vl=min(VLEN/SEW,rs1), rd=vl**

Resulting machine vector length setting      Requested application vector length

**vttype** parameters (SEW,LMUL,EDIV) encoded as immediate in instruction

Instruction encoding



Usually use immediate form, **vsetvli**, to set **vttype** parameters.

The register version **vsetvl** is usually used only for context save/restore

# vsetvl{i} Instruction

- The first scalar register argument, *rs1*, is the requested application vector length (AVL)
  - The type argument (either immediate or second register *rs2*) indicates how the vector registers should be configured
- Configuration includes size of each element, SEW, and LMUL value
- The vector length is set to the *minimum* of requested AVL and the maximum supported vector length (VLMAX) in the new configuration
  - $VLMAX = LMUL * VLEN / SEW$
  - $vl = \min(AVL, VLMAX)$
- The value placed in **vl** is also written to the scalar destination register *rd*

# Stripmined vector Memory Copy Example

```
# void *memcpy(void* dest, const void* src, size_t n)
# a0=dest, a1=src, a2=n
#
Set configuration, calculate vector strip length,
Unit-stride vector load elements (bytes),
Unit-stride vector store elements (bytes)
memcpy:
    mv a3, a0 # Copy destination
loop:
    vsetvli t0, a2, e8 # Vectors of 8b
    vle.v v0, (a1)      # Load bytes
    add a1, a1, t0      # Bump pointer
    sub a2, a2, t0      # Decrement count
    vse.v v0, (a3)      # Store bytes
    add a3, a3, t0      # Bump pointer
    bnez a2, loop       # Any more?
    ret                # Return
```

*Same binary machine code can run on machines with any VLEN!*



# Mixed Precision Operations Support

- The vector ISA is designed to support mixed-width operations without requiring explicit additional rearrangement instructions.
- The recommended software strategy is to modify vtype dynamically to keep SEW/LMUL constant (and hence VLMAX constant) when operating on vectors of different precision values.
- The following example shows four different packed element widths (8b, 16b, 32b, 64b) in a VLEN=128b implementation. The vector register grouping factor (LMUL) is increased by the relative element size such that each group can hold the same number of vector elements (VLMAX=8 in this example) to simplify stripmining code.

Example VLEN=128b, with SEW/LMUL=16

[illegible]

# Mask Register

- Nearly all operations can be optionally under a mask (or predicate) held in vector register v0
- A single vm bit in instruction encoding selects whether unmasked or under control of v0
- A vector mask occupies only one vector register regardless of SEW and LMUL.
- Each element is allocated a single mask bit in a mask vector register.
- The mask bit for element i is located in bit i of the mask register, independent of SEW or LMUL.

VLEN=32b

Byte      3    2    1    0  
LMUL=1, SEW=8b

          3    2    1    0    Element  
          [03][02][01][00] Mask bit position in decimal

LMUL=2, SEW=16b


          1            0  
          [01]        [00]  
          3            2  
          [03]        [02]

LMUL=4, SEW=32b

                  0  
                  [00]  
                  1  
                  [01]  
                  2  
                  [02]  
                  3  
                  [03]


# Vector Load Instructions

unit-stride




strided

Scalar stride (bytes)



indexed


Vector of offsets (bytes)




Scalar base address

# Vector store Instructions


unit-stride



strided



indexed



## Vector Unit Stride Load Stores

```
# vd destination, rs1 base address, vm is mask encoding (v0.t or <missing>)  
vle.v          vd, (rs1), vm      # SEW  
# vs3 store data, rs1 base address, vm is mask encoding (v0.t or <missing>)  
vse.v          vs3, (rs1), vm     # SEW store
```

## Vector Strided Load Stores

```
# vd destination, rs1 base address, rs2 byte stride  
vlse.v         vd, (rs1), rs2, vm # SEW  
# vs3 store data, rs1 base address, rs2 byte stride  
vsse.v         vs3, (rs1), rs2, vm # SEW
```

## Vector Indexed Load Stores

```
# vd destination, rs1 base address, vs2 indices  
vlxe.v         vd, (rs1), vs2, vm # SEW  
# vs3 store data, rs1 base address, vs2 indices  
vsxe.v         vs3, (rs1), vs2, vm # SEW
```

# Week-12 16 Oct 2020

1. RISC-V ISA Vector Extension  
Instruction Set {Chapters 7 – 17}
2. Further Course of Action

## ARITHMETIC AND LOGICAL VECTOR INSTRUCTIONS

ADD SUB MUL DIV FMA CMP CVT LOGIC CLIP SHIFT REDUCTION

INTEGER

FLOATING POINT

FP

WIDENING / NARROWING

W/N

VECTOR - VECTOR






VECTOR - SCALAR

V-S

VECTOR - IMMEDIATE

1. There is NO SPERERATE *FP Vector Register file*. Both Integer and FP values are stored in the same Vector Register files.
2. Nearly all operations can be optionally under a mask (or predicate) held in vector register v0.
3. A single vm bit in instruction encoding selects whether unmasked or under control of v0

# Vector Arithmetic Instruction Encodings





# Vector Add Instructions

# Integer adds.

```
vadd.vv vd, vs2, vs1, vm    # Vector-vector  
vadd.vx vd, vs2, rs1, vm    # vector-scalar  
vadd.vi vd, vs2, imm, vm    # vector-immediate
```

# Integer subtract

```
vsub.vv vd, vs2, vs1, vm    # Vector-vector  
vsub.vx vd, vs2, rs1, vm    # vector-scalar
```

# Integer reverse subtract

```
vrsub.vx vd, vs2, rs1, vm    #  $vd[i] = rs1 - vs2[i]$   
vrsub.vi vd, vs2, imm, vm    #  $vd[i] = imm - vs2[i]$ 
```

# Floating-point add

```
vfadd.vv vd, vs2, vs1, vm    # Vector-vector  
vfadd.vf vd, vs2, rs1, vm    # vector-scalar
```

# Floating-point subtract

```
vfsb.vv vd, vs2, vs1, vm    # Vector-vector  
vfsb.vf vd, vs2, rs1, vm    # Vector-scalar  $vd[i] = vs2[i] - f[rs1]$   
vfrsub.vf vd, vs2, rs1, vm    # Scalar-vector  $vd[i] = f[rs1] - vs2[i]$ 
```

# Widening Integer Add Instructions

```
# Widening unsigned integer add/subtract, 2*SEW = SEW +/- SEW
vwaddu.vv vd, vs2, vs1, vm # vector-vector
vwaddu.vx vd, vs2, rs1, vm # vector-scalar
vwsubu.vv vd, vs2, vs1, vm # vector-vector
vwsubu.vx vd, vs2, rs1, vm # vector-scalar
```

```
# Widening signed integer add/subtract, 2*SEW = SEW +/- SEW
vwadd.vv vd, vs2, vs1, vm # vector-vector
vwadd.vx vd, vs2, rs1, vm # vector-scalar
vwsub.vv vd, vs2, vs1, vm # vector-vector
vwsub.vx vd, vs2, rs1, vm # vector-scalar
```

```
# Widening unsigned integer add/subtract, 2*SEW = 2*SEW +/- SEW
vwaddu.wv vd, vs2, vs1, vm # vector-vector
vwaddu.wx vd, vs2, rs1, vm # vector-scalar
vwsubu.wv vd, vs2, vs1, vm # vector-vector
vwsubu.wx vd, vs2, rs1, vm # vector-scalar
```

```
# Widening signed integer add/subtract, 2*SEW = 2*SEW +/- SEW
vwadd.wv vd, vs2, vs1, vm # vector-vector
vwadd.wx vd, vs2, rs1, vm # vector-scalar
vwsub.wv vd, vs2, vs1, vm # vector-vector
vwsub.wx vd, vs2, rs1, vm # vector-scalar
```

# Widening FP Fused Mult-Acc

# FP widening multiply-accumulate, overwrites addend

`vfwmacc.vv vd, vs1, vs2, vm`    #  $vd[i] = +(vs1[i] * vs2[i]) + vd[i]$

`vfwmacc.vf vd, rs1, vs2, vm`    #  $vd[i] = +(f[rs1] * vs2[i]) + vd[i]$

# FP widening negate-(multiply-accumulate), overwrites addend

`vfwnmacc.vv vd, vs1, vs2, vm`    #  $vd[i] = -(vs1[i] * vs2[i]) - vd[i]$

`vfwnmacc.vf vd, rs1, vs2, vm`    #  $vd[i] = -(f[rs1] * vs2[i]) - vd[i]$

# FP widening multiply-subtract-accumulator, overwrites addend

`vfwsac.vv vd, vs1, vs2, vm`    #  $vd[i] = +(vs1[i] * vs2[i]) - vd[i]$

`vfwsac.vf vd, rs1, vs2, vm`    #  $vd[i] = +(f[rs1] * vs2[i]) - vd[i]$

# FP widening negate-(multiply-subtract-accumulator), overwrites addend

`vfwnmsac.vv vd, vs1, vs2, vm`    #  $vd[i] = -(vs1[i] * vs2[i]) + vd[i]$

`vfwnmsac.vf vd, rs1, vs2, vm`    #  $vd[i] = -(f[rs1] * vs2[i]) + vd[i]$

# Vector Reduction

- Take a Vector Register group of elements and a scalar held in element 0 of a vector register(vs1[0]), and perform a reduction operation using some binary operator to produce a scalar result in element 0 of a vector register(vd[0]).
- Supports FP, Widening, Ordered and Un-ordered sums

```
# Simple reductions, where [*] denotes all active elements:
vredsum.vs  vd, vs2, vs1, vm  # vd[0] =  sum( vs1[0] , vs2[*] )
vredmaxu.vs vd, vs2, vs1, vm  # vd[0] = maxu( vs1[0] , vs2[*] )
vredmax.vs  vd, vs2, vs1, vm  # vd[0] =  max( vs1[0] , vs2[*] )
vredminu.vs vd, vs2, vs1, vm  # vd[0] = minu( vs1[0] , vs2[*] )
vredmin.vs  vd, vs2, vs1, vm  # vd[0] =  min( vs1[0] , vs2[*] )
vredand.vs  vd, vs2, vs1, vm  # vd[0] =  and( vs1[0] , vs2[*] )
vredor.vs   vd, vs2, vs1, vm  # vd[0] =   or( vs1[0] , vs2[*] )
vredxor.vs  vd, vs2, vs1, vm  # vd[0] =  xor( vs1[0] , vs2[*] )
```

```
# Simple reductions.
```

```
vfwredosum.vs vd, vs2, vs1, vm # Ordered sum
```

```
vfwredsum.vs  vd, vs2, vs1, vm # Unordered sum
```

# Recent Developments in RISCV-V Ext

- Andes Group develops Vector based Core for RISCV -V Specifications (First Company to do so) - May 26, 2020
- GNU-tool chain for Vector Spec v0.8 available on GitHub
- RISC-V V likely to finalize in Ver1.0

# Future Course of Action(s)

- Micro-architecture design
  - Very basic outline provided in Patterson Hennessy text(5th Edition)
  - Krste Asanovac's PhD thesis (1998) discusses microarchitectural design in slightly more detail (reg-file, memory design (w/o cache),
  - Couple of papers such as Hwacha, Ara (Open Source, not gone thru yet)
- High level design for integrating Vector processor into RISCV core
  - Interface between main core and vector core
  - How closely coupled ? On which bus should the Vector Unit be present Main bus or peripheral(AXI) bus?
- Design of FPU (16,32)
  - I'm unfamiliar
  - IP-Core available to integrate OR else, design from scratch?
- ISA-Simulator coding required ? yes
- Resources on FPGA . Require Part number details

# Ordered and Unordered Sums

- **Ordered Sum** Must sum the values in element order, starting with the scalar in `vs1[0]`--that is, it performs the computation:  $((vs1[0] + vs2[0]) + vs2[1]) + \dots + vs2[vl-1]$ , where each addition operates identically to the scalar floating-point instructions in terms of raising exception flags and generating or propagating special values. Code good for compiler to auto-vectorize.
- **Unordered Sum** provides freedom to the implementation to perform reduction – Binary Tree of all vector operands followed by  $\log_2(n)$  steps of addition. This is faster

# Vector Permutation Instructions

- Provided to move elements around within the vector registers.
- **Integer Scalar Move Instruction** transfers a single value between a scalar x register and element 0 of a vector register.
- **Vector Slide Instructions** move vector elements up or down in a vector register group
- **Vector Register Gather Instruction** reads elements from a first source vector register group at locations given by a second source vector register group.
- **Vector Compress Instruction** allows elements selected by a vector mask register from a source vector register group to be packed into contiguous elements at the start of the destination vector register group.
- **Whole Vector Register Move Instruction** copies all VLEN bits in rs -->rd